

# Differential Gene Expression

RNA-seq Data Analysis Course, EBI, April 2020

Dr. Alexey Larionov

Department of Medical Genetics  
Cambridge University, UK

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

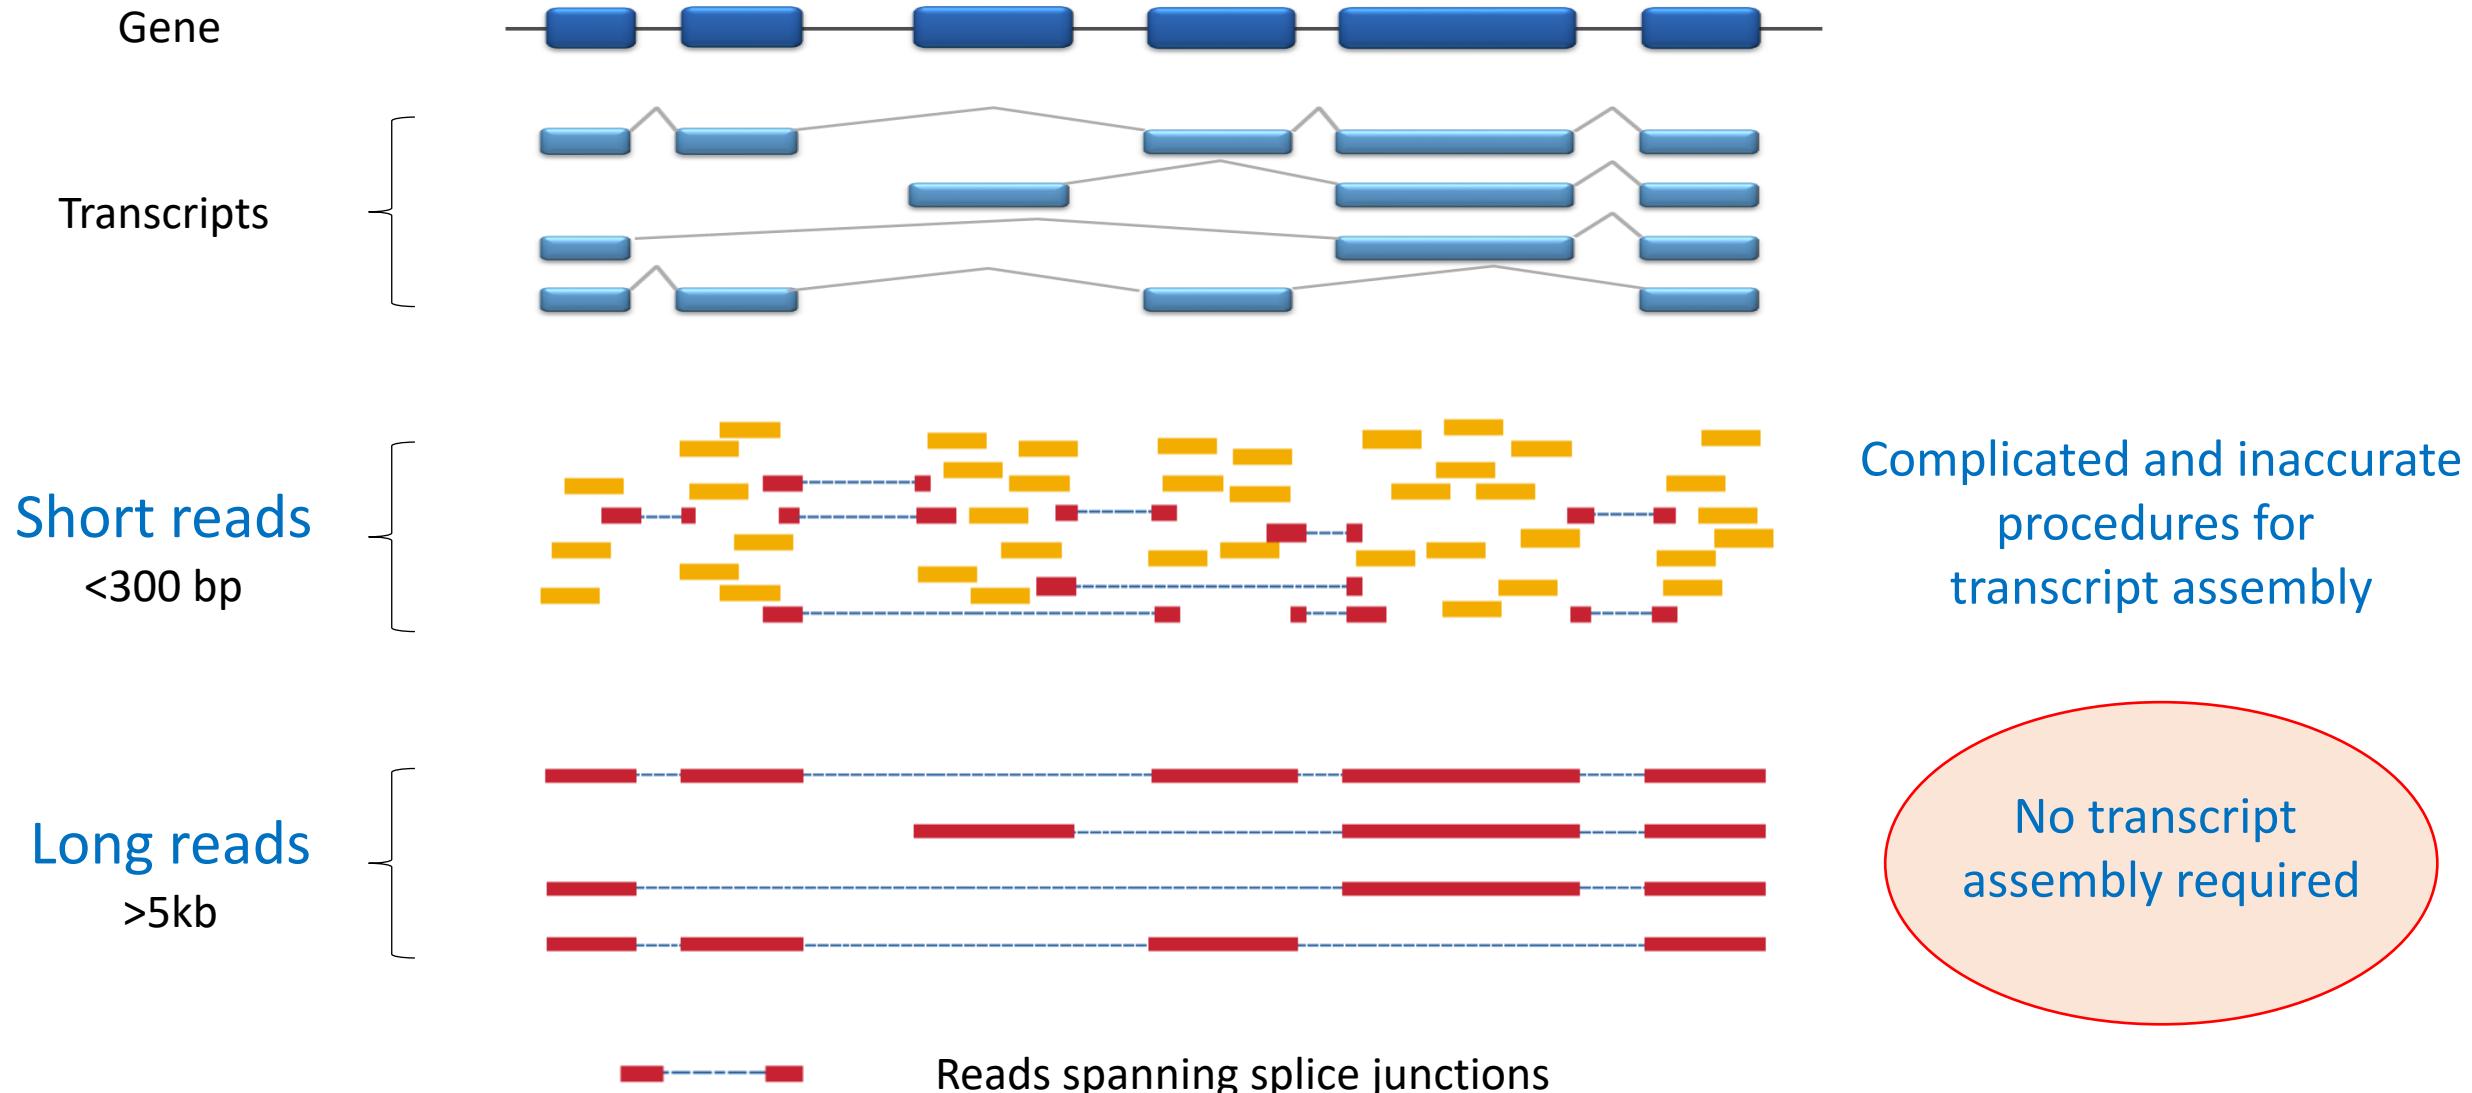
## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

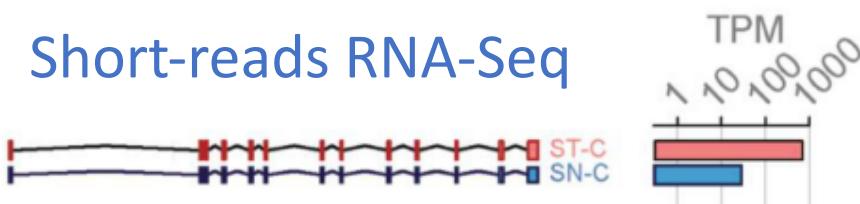
Comparison of methods  
Comparison of results

# Short and long-reads in RNA-seq



# Short and long-reads in RNA-seq

## Short-reads RNA-Seq



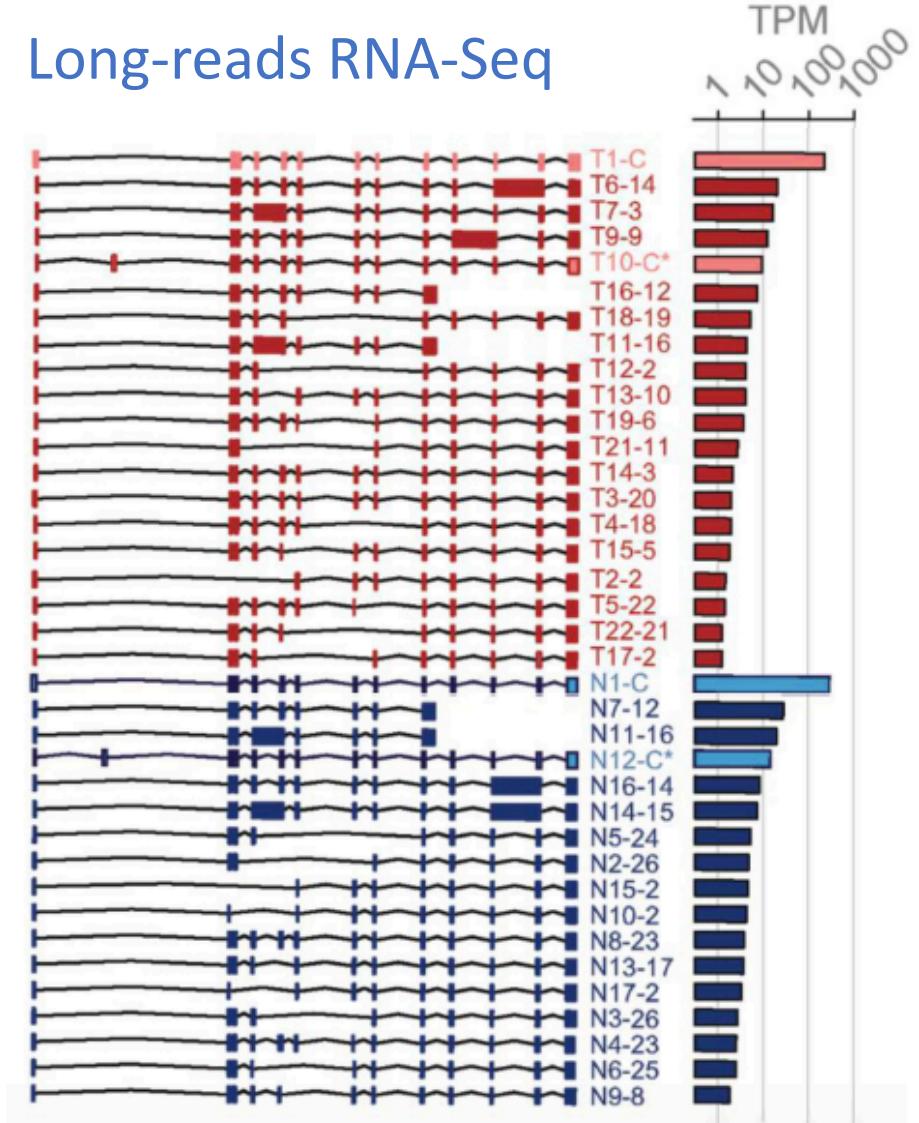
ANXA1 gene expression measurement  
Sample A vs Sample B

... full-length RNA-sequencing ... revealed a ~5-fold higher number of transcript isoforms than previously detected

Mays *et al* 2019

Tools for short reads are not good for long reads  
Tools for long reads are yet emerging  
(sometime platform-specific, e.g. ToFU for PacBio)

## Long-reads RNA-Seq



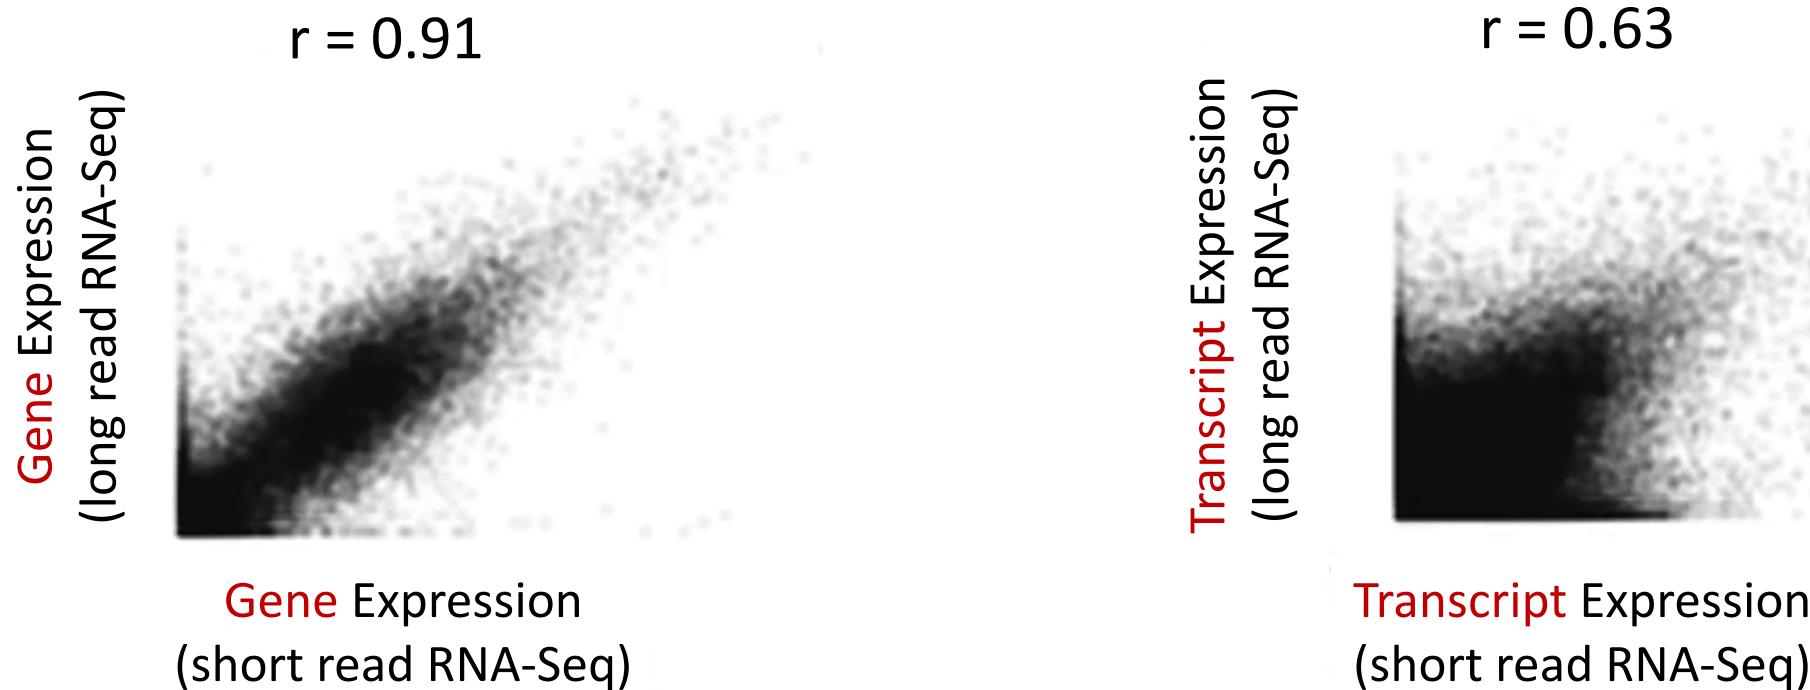
Mays et al 2019 Single-Molecule Real-Time (SMRT) Full-Length RNA-Sequencing Reveals Novel and Distinct mRNA Isoforms in Human Bone Marrow Cell Subpopulations Genes. 2019, 10, 253

ToFU: Gordon et al. 2015 Widespread Polycistronic Transcripts in Fungi Revealed by Single-Molecule mRNA Sequencing. PLoS ONE 10: e0132628

MISO: Katz et al 2010 Analysis and design of RNA sequencing experiments for identifying isoform regulation Nature Methods 7 1009

# Short and long-reads in RNA-seq

Gene expressions correlate well between short- and long – reads  
Transcript expressions do not correlate well



# Differential **Gene** Expression

RNA-seq Data Analysis Course, EBI, April 2020

Dr. Alexey Larionov

Department of Medical Genetics  
Cambridge University, UK

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts

### Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches

Problems with RNA-seq counts

Overdispersion: Negative Binomial Distribution

## Counts

Overview and software

Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq

GTF file format

## Design formula

A simple design

Accounting for covariates e.g. batch effect

Advanced designs: interactions

## Data import

Summarized Experiment and DGEList

Data import packages and functions

## Statistics in more details

Normalizing by library size

Exploring the source data

- Variance-stabilizing transformations
- PCA and Hierarchical clustering

Dispersion estimates

- Need for “borrowing” data
- Empirical adaptive estimates

Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods

Comparison of results

# RNA-Seq gene expression analysis

## Alignment to genome

with gapped aligners  
STAR, (r)subreads-align,  
minimap2, GMAP

## Transcripts assembly

from genome-alignment  
StringTie, Cufflinks, ToFU

## Reads

FastQC, Trimming adaptors  
Removing poor bases and reads  
Cutadapt, Trimmomatic

**De-novo transcriptome assembly**  
becoming possible with long reads  
...

## Alignment to transcriptome

with un-gapped aligners  
Bowtie

## Alignment-free quantification

Salmon , Kallisto, Sailfish

## Counting reads

(r)subreads-featureCounts  
RSEM, HTSeq-count

## Transcripts/genes counts

## Differential expression

DESeq2, edgeR, limma-voom, MISO

## Functional analysis

GO, Pathways, GSEA

## Classification

Signatures, linear models, co-clustering  
Machine learning: Random Forrest, ANN

# RNA-Seq gene expression analysis

## Alignment to genome

with gapped aligners  
STAR, (r)subreads-align,  
minimap2, GMAP

## Transcripts assembly

from genome-alignment  
StringTie, Cufflinks, ToFU

## Reads

FastQC, Trimming adaptors  
Removing poor bases and reads  
Cutadapt, Trimmomatic

**De-novo transcriptome assembly**  
becoming possible with long reads  
...

## Alignment to transcriptome

with un-gapped aligners  
Bowtie

## Counting reads

(r)subreads-featureCounts  
RSEM, HTSeq-count

## Alignment-free quantification

Salmon , Kallisto, Sailfish

## Transcripts/genes counts

### Differential expression

DESeq2, edgeR, limma-voom, MISO

This lecture focuses only on **genes** differential expression although statistical principles could be extended to other types of counts

This slide focuses on options for short reads. Tools and practices for long reads are yet emerging

## Functional analysis

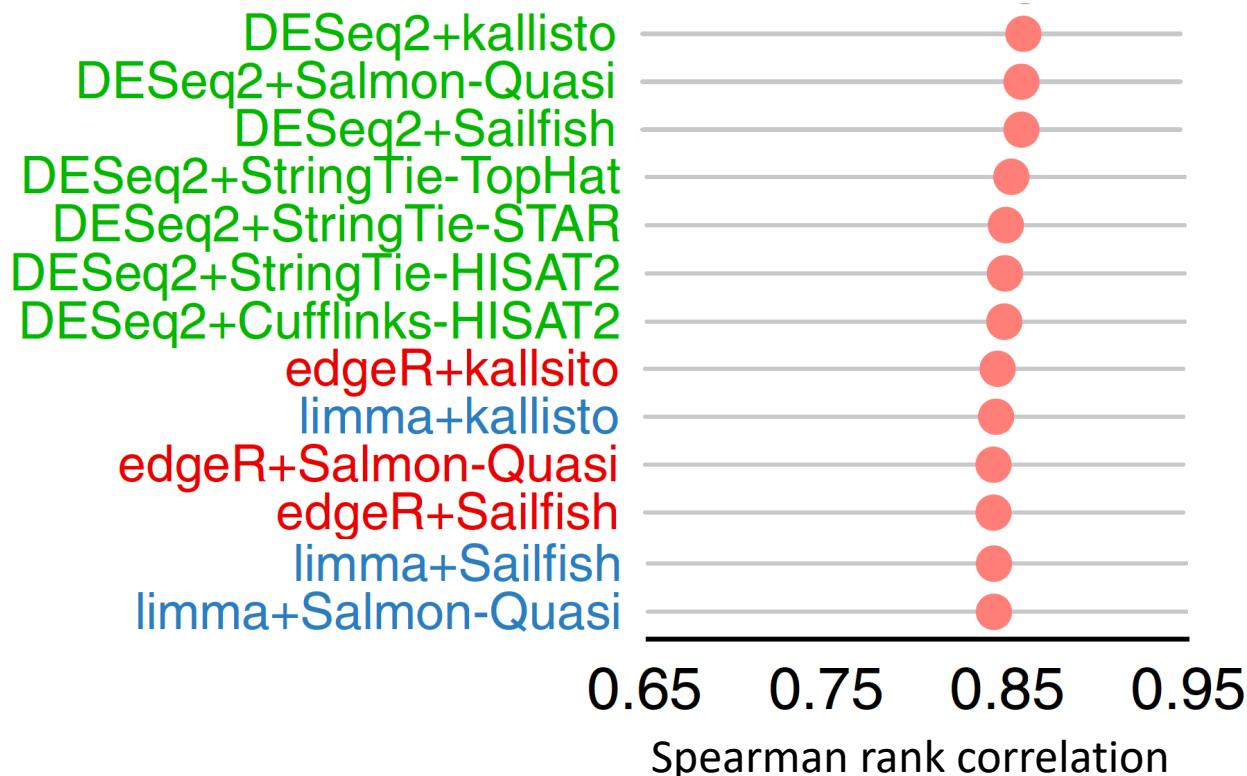
GO, Pathways, GSEA

## Classification

Signatures, linear models, co-clustering  
Machine learning: Random Forrest, ANN

# Comparison of tools for differential expression analysis in short-read RNA-Seq

## RNA-Seq vs qRT-PCR



1,001 genes were measured in two samples by RNA-Seq and by qRT-PCR  
Advantage of DESeq2 over edgeR/limma was even stronger in other comparisons

# Differential Gene Expression

## Short reads RNA-seq Analysis with DESeq2 and edgeR

Dr. Alexey Larionov

Department of Medical Genetics  
Cambridge University, UK

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

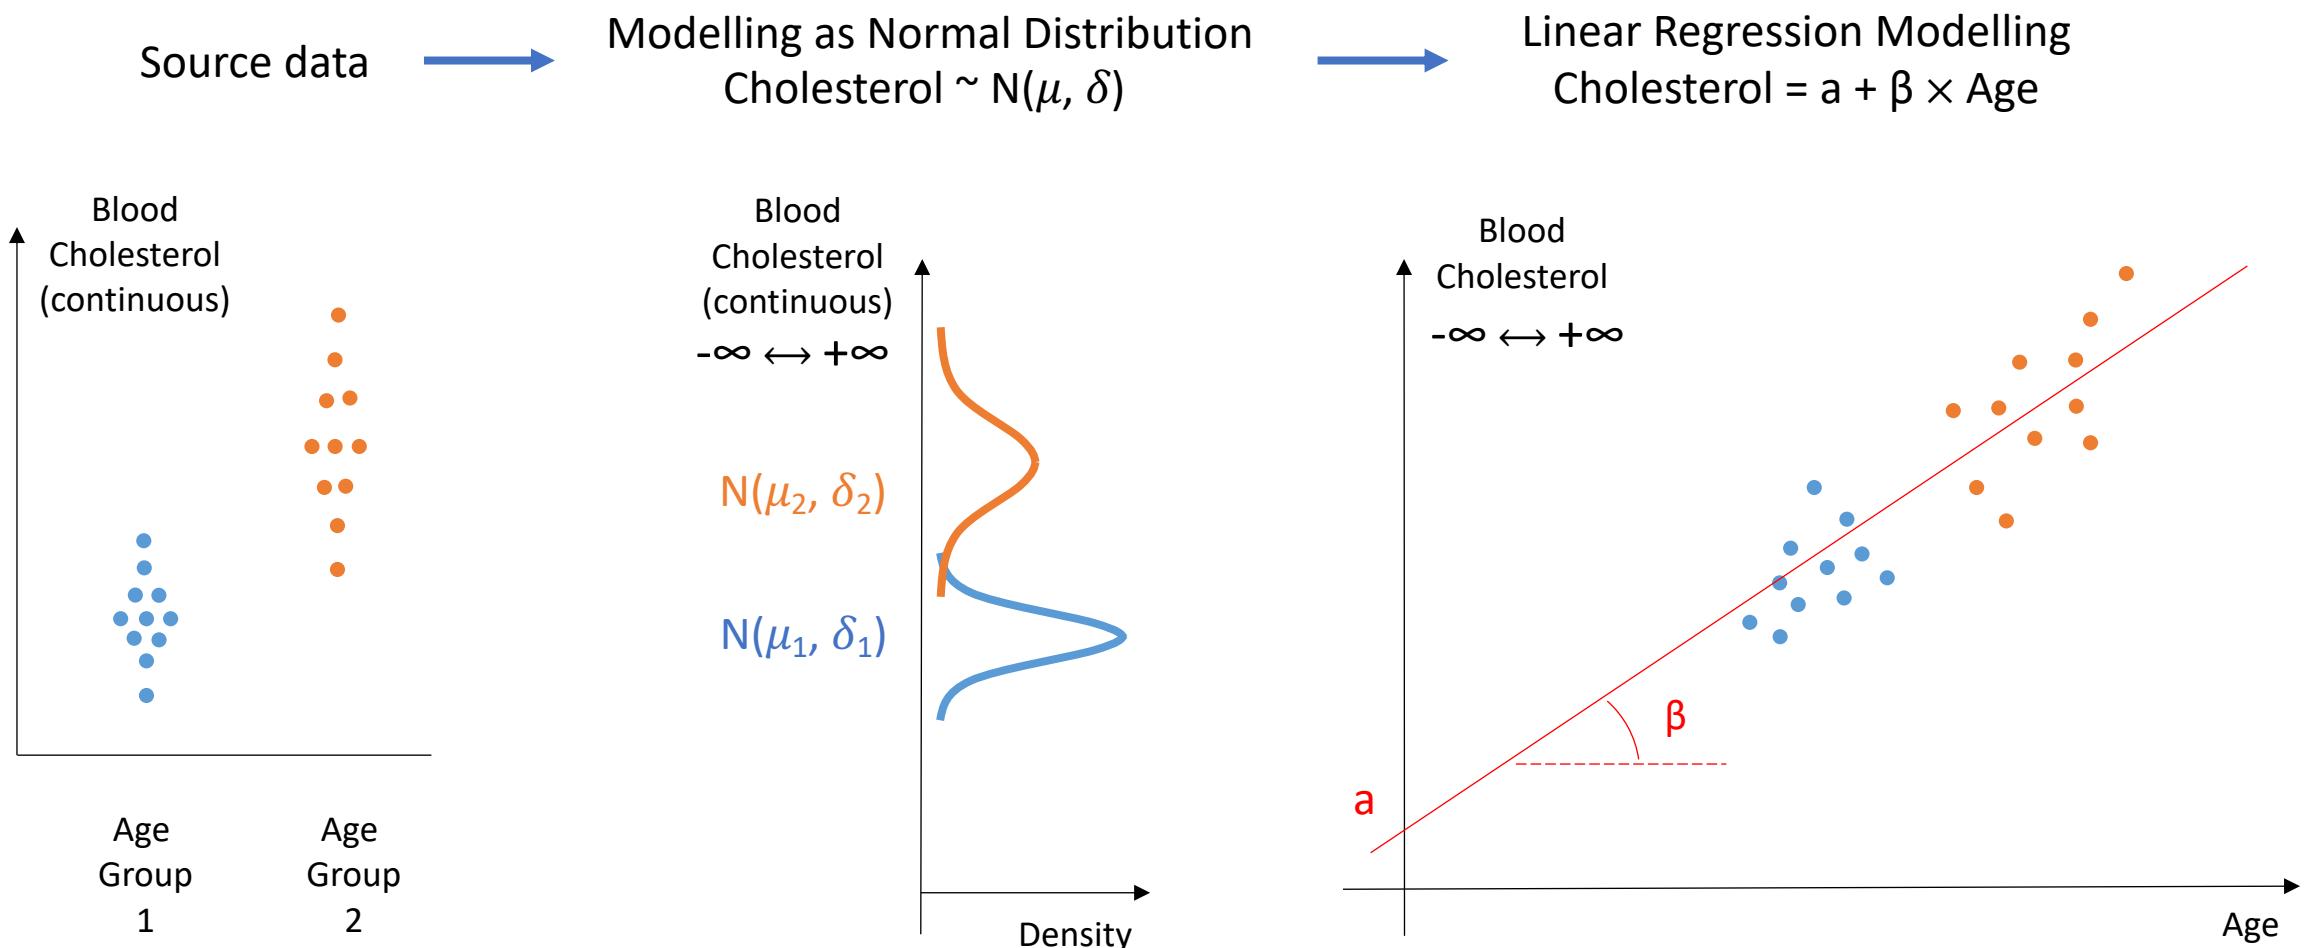
## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

# Recap of statistical approaches for detecting difference between groups



Visual assessment suggests a trend

Formal significance assessment  
e.g. Student's test  
 $\mu_1 \neq \mu_2$  at  $p < 0.05$

Significance of  $\beta \neq 0$ ;  
Convenient to add covariates (confounders), e.g.:  
 $\text{Cholesterol} = a + \beta_1 \times \text{Weight} + \beta_0 \times \text{Age}$

# Why can't we apply this framework for Differential Gene Expression Analysis ?

## Problems

- 1) Raw counts in each sample depend on library size (depth of sequencing)
- 2) Low counts do not obey the “Normal” bell-shape distribution because they can’t go below zero
- 3) The counts are discrete, which is better modelled by a discrete distribution
- 4) Small number of samples does not allow accurate estimation of dispersion (variance)
- 5) Testing for many genes at a time

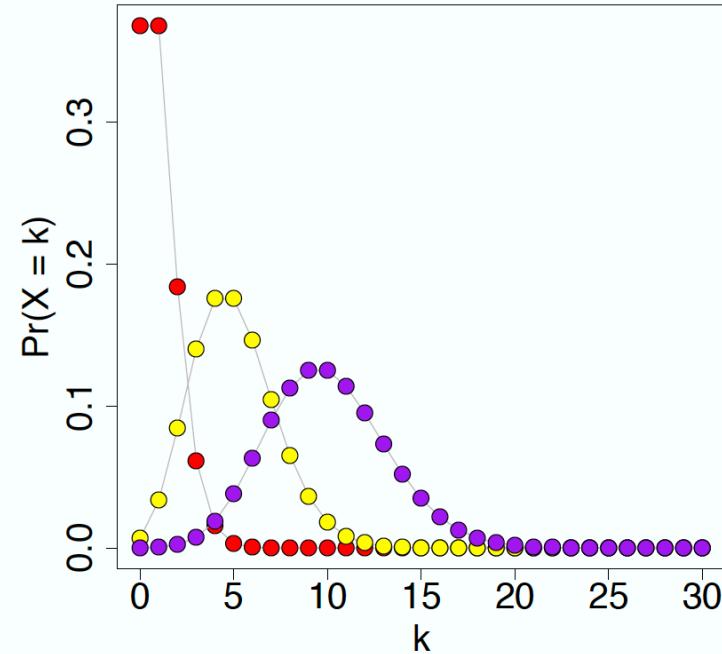
## Solutions

- 1) Normalizing raw counts by the library size (discussed in later)
- 2) and 3) Choosing an appropriate discrete distribution
- 4) “Borrowing” data between genes for estimation of dispersion (discussed later)
- 5) Multiple testing correction (typically FDR)

# Poisson distribution

Distribution of random independent events happening at a certain **mean** rate.  
Mathematically, the dispersion(variance) is equal to the mean.

$$P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$



By design describes the random sampling of molecules from a solution with given concentration.

Exactly matches the counts distribution in the technical replicas of RNA-seq:  
e.g. sequencing of several aliquots from the same library.

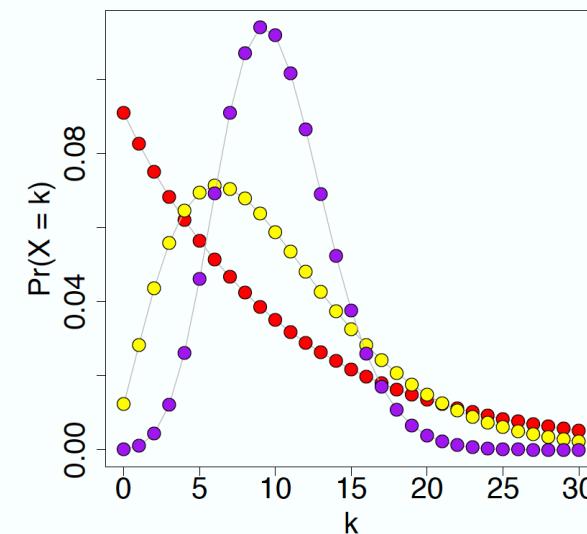
# Overdispersion

Total variance	=	Technical variance i.e. between replicas within library, described by Poisson distribution	+	Additional variance e.g. between dishes of the same cell line or different tumors of the same type
----------------	---	--	---	--

## Negative Binomial Distribution

Number of independent attempts until a certain number of successes  
Mathematically, allows dispersion larger than mean

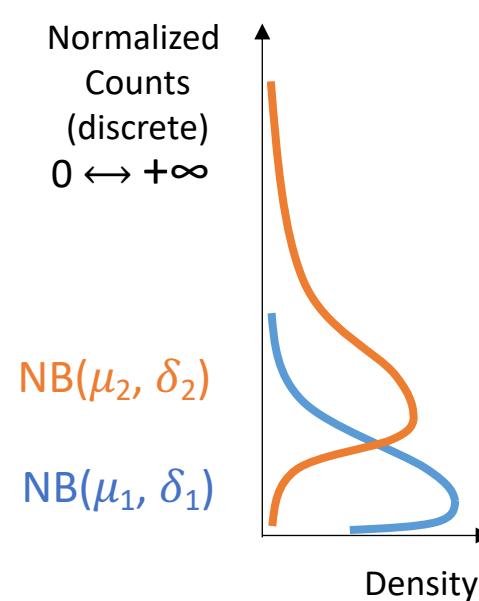
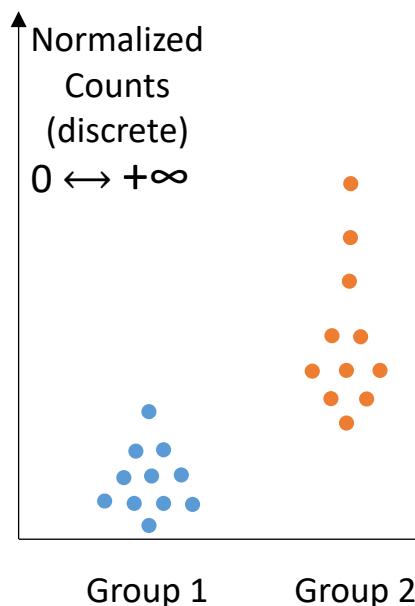
$$P(X = k) = \binom{k + r - 1}{k} \cdot (1 - p)^r p^k$$



“Similar” to Poisson: discrete and non-negative. However, unlike to Poisson allows to model the overdispersion.  
Successfully used to model real-life RNA-seq data (details about the dispersion assessment will be discussed later).

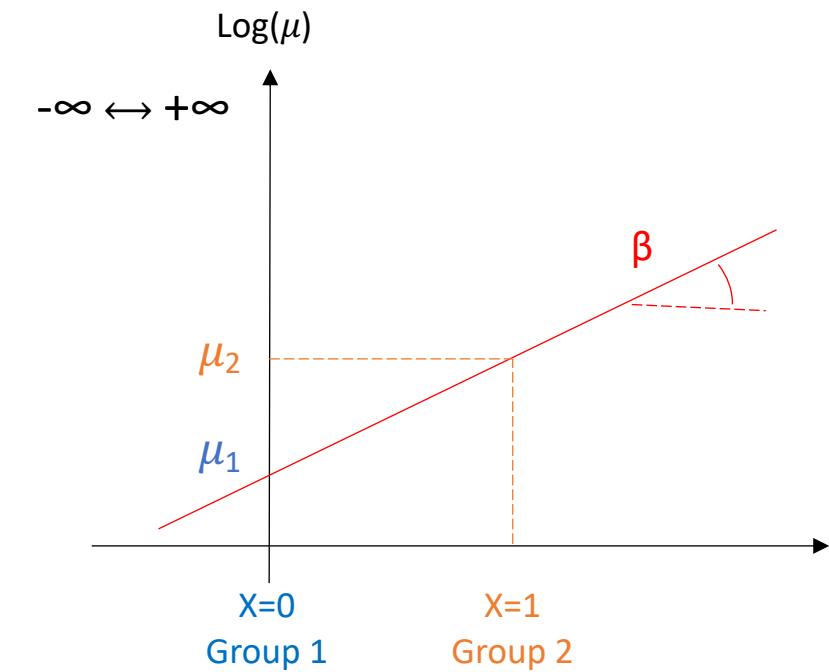
# Overview of statistical approaches to Differential Genes Expression analysis

Source data  
Normalized by library size → Modelling as Negative Binomial  
Counts  $\sim \text{NB}(\mu, \delta)$  → Log-link → Generalized Linear Modeling  
 $\text{Log}(\mu) = a + \beta \times \text{Group}$



Visual assessment suggests a trend

“Borrow” data for  
“adaptive” variance evaluation:  
for low number of cases  
 $\delta(\text{gene})$  shrinks to  
 $\delta(\text{all genes with similar expression})$



Significance of  $\beta \neq 0$   
Convenient adding covariates (confounders):  
Counts =  $a + \beta_1 \times \text{Batch} + \beta_0 \times \text{Group}$

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

# RNA-Seq gene expression analysis

## Alignment to genome

with gapped aligners  
STAR, (r)subreads-align,  
minimap2, GMAP

## Transcripts assembly

from genome-alignment  
StringTie, Cufflinks, ToFU

## Alignment to transcriptome

with un-gapped aligners  
Bowtie

## Counting reads

(r)subreads-featureCounts  
RSEM, HTSeq-count

## Reads

FastQC, Trimming adaptors  
Removing poor bases and reads  
Cutadapt, Trimmomatic

## De-novo transcriptome assembly

becoming possible with long reads  
...

## Alignment-free quantification

Salmon , Kallisto, Sailfish

## Transcripts/genes counts

## Differential expression

DESeq2, edgeR, limma-voom, MISO

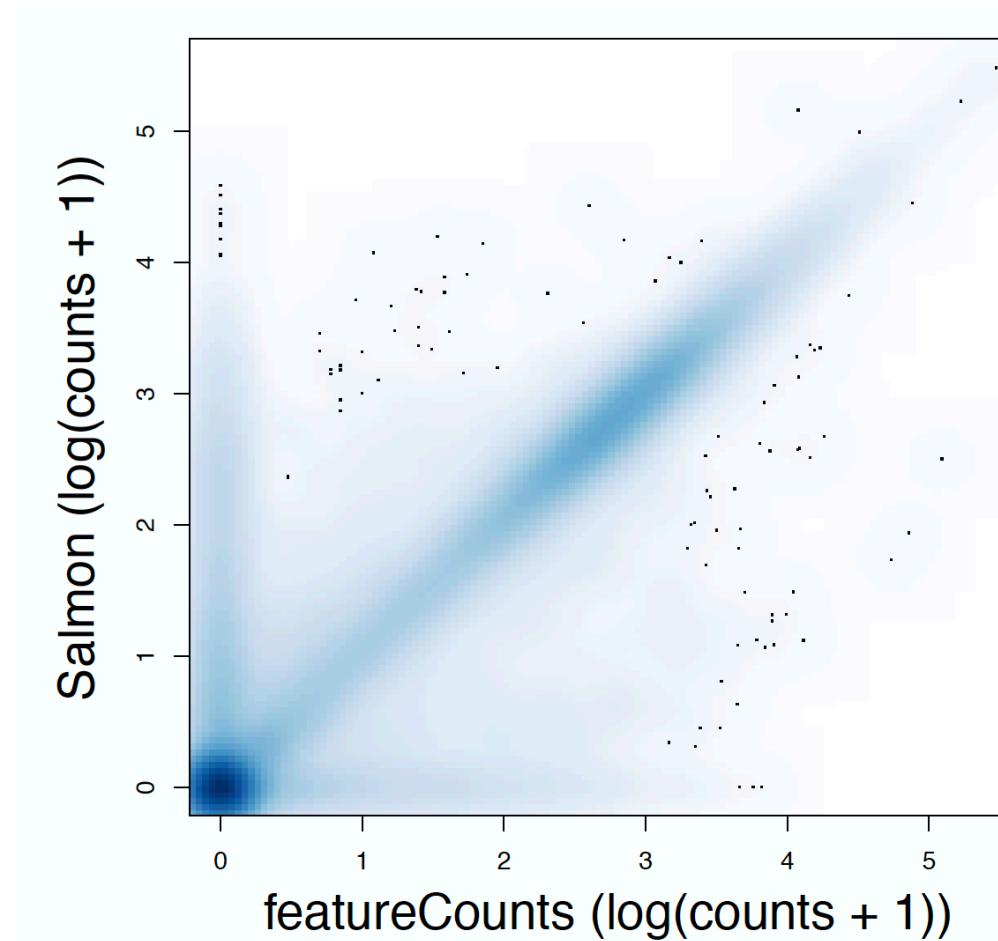
## Functional analysis

GO, Pathways, GSEA

## Classification

Signatures, linear models, co-clustering  
Machine learning: Random Forrest, ANN

Gene-level counts are reasonably close but not identical  
between different approaches



# Caveats in getting raw counts ...

For long-reads:

count of transcripts is reasonably natural  
read ~ transcript (within reason...)

For short-reads:

count reads intersecting with features...

How to intersect ?

Options in **htseq** library :

<http://htseq.readthedocs.io/en/master/count.html>

What **features** to use ?

Transcripts, genes or exons ...

	union	intersection _strict	intersection _nonempty
A single green horizontal bar labeled "read" overlaps a purple horizontal bar labeled "gene_A".	gene_A	gene_A	gene_A
A green horizontal bar labeled "read" starts within a purple horizontal bar labeled "gene_A".	gene_A	no_feature	gene_A
Two overlapping green horizontal bars labeled "read" are shown on a purple horizontal bar labeled "gene_A".	gene_A	no_feature	gene_A
Two overlapping green horizontal bars labeled "read" are on a purple horizontal bar labeled "gene_A", and one green horizontal bar labeled "read" is on a blue horizontal bar labeled "gene_B".	gene_A	gene_A	gene_A
A green horizontal bar labeled "read" starts within a purple horizontal bar labeled "gene_A", and another green horizontal bar labeled "read" starts within a blue horizontal bar labeled "gene_B".	gene_A	gene_A	gene_A
A green horizontal bar labeled "read" overlaps both a purple horizontal bar labeled "gene_A" and a blue horizontal bar labeled "gene_B".	ambiguous	gene_A	gene_A
A green horizontal bar labeled "read" starts within a purple horizontal bar labeled "gene_A" and ends within a blue horizontal bar labeled "gene_B".	ambiguous	ambiguous	ambiguous

# Units for read counts : raw counts + RPKM, FPKM, TPM

## Historically used units

**RPKM** = Reads Per Kilobase of the feature (gene) per Million

**FPKM** = Fragments Per Kilobase of the feature (gene) per Million

$$R(F)PKM = \frac{\text{Number of reads (fragments) mapped for transcript} \times 10^3 \times 10^6}{\text{Transcript length} \times \text{Number of reads (fragments) mapped in sample}}$$

## Currently recommended unit

**TPM** = Transcripts Per Million

$$TPM = \frac{\text{Number of fragments mapped for transcript} \times \text{Average fragment length} \times 10^6}{\text{Length of transcript} \times \text{Number of transcripts in sample}}$$

Accounts for transcript and library sizes

# Units for read counts : raw counts + RPKM, FPKM, TPM

Raw counts are needed  
for Differential Expression  
analysis !

## Historically used units

**RPKM** = Reads Per Kilobase of the feature (gene) per Million

**FPKM** = Fragments Per Kilobase of the feature (gene) per Million

$$R(F)PKM = \frac{\text{Number of reads (fragments) mapped for transcript} \times 10^3 \times 10^6}{\text{Transcript length} \times \text{Number of reads (fragments) mapped in sample}}$$

## Currently recommended unit

**TPM** = Transcripts Per Million

$$TPM = \frac{\text{Number of fragments mapped for transcript} \times \text{Average fragment length} \times 10^6}{\text{Length of transcript} \times \text{Number of transcripts in sample}}$$

Accounts for transcript and library sizes

# Different tools provide counts in different formats

## HTSeq

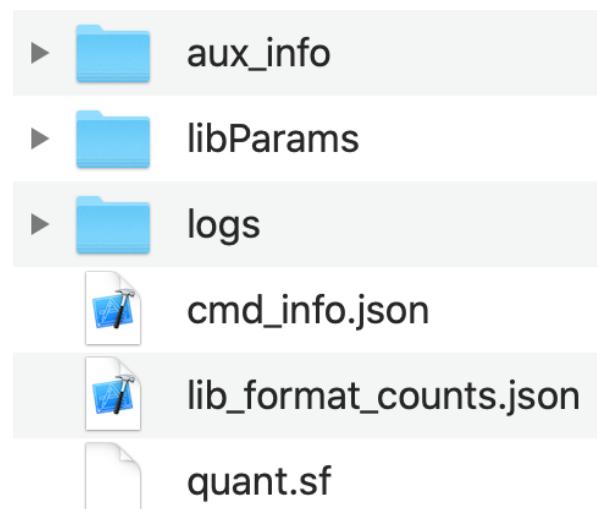
Separate text file for each sample

A screenshot of a terminal window titled "example.htseq.counts". The window displays a list of gene IDs and their corresponding read counts. The data includes several genes with non-zero counts and a large number of genes with zero counts, such as "no\_feature" and "ambiguous".

ENSG00000000003.13	2569
ENSG00000000005.5	1
ENSG00000000419.11	3180
ENSG00000000457.12	3332
ENSG00000000460.15	1621
ENSG00000000938.11	530
ENSG00000000971.14	7282
ENSG00000001036.12	3312
ENSG00000001084.9	2642
ENSG00000001167.13	3322
...	
ENSGR0000270726.4	0
ENSGR0000275287.3	0
ENSGR0000276543.3	0
ENSGR0000277120.3	0
ENSGR0000280767.1	0
ENSGR0000281849.1	0
__no_feature	3069305
__ambiguous	3368739
__too_low_aQual	0
__not_aligned	0
__alignment_not_unique	23748640

## Salmon

Separate text file for each sample  
in a folder with meta-information



## (r)subreads featureCount

Matrix of counts with  
Genes in rows and  
Samples in columns

A screenshot of a terminal window titled "quant.sf". The window displays a table of transcript-level gene expression metrics. The columns are labeled "Name", "Length", "EffectiveLength", "TPM", and "NumReads". The data shows various transcripts with their respective lengths, effective lengths, TPM values, and numbers of reads.

Name	Length	EffectiveLength	TPM	NumReads
ENST00000456328.2	1657	1675.004	0.843814	17.896
ENST00000450305.2	632	477.000	0.000000	0.000
ENST00000488147.1	1351	1332.253	3.955642	66.727
ENST00000619216.1	68	11.000	0.000000	0.000
ENST00000473358.1	712	557.000	0.000000	0.000
ENST00000469289.1	535	381.000	0.000000	0.000
ENST00000607096.1	138	27.000	0.000000	0.000
ENST00000417324.1	1187	1032.000	0.000000	0.000
ENST00000461467.1	590	436.000	0.000000	0.000
ENST00000606857.1	840	685.000	0.000000	0.000
ENST00000642116.1	1414	1259.000	0.000000	0.000
...				

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

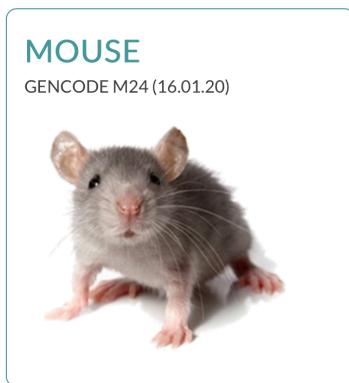
## Visualizing results

MA- and Volcano plots

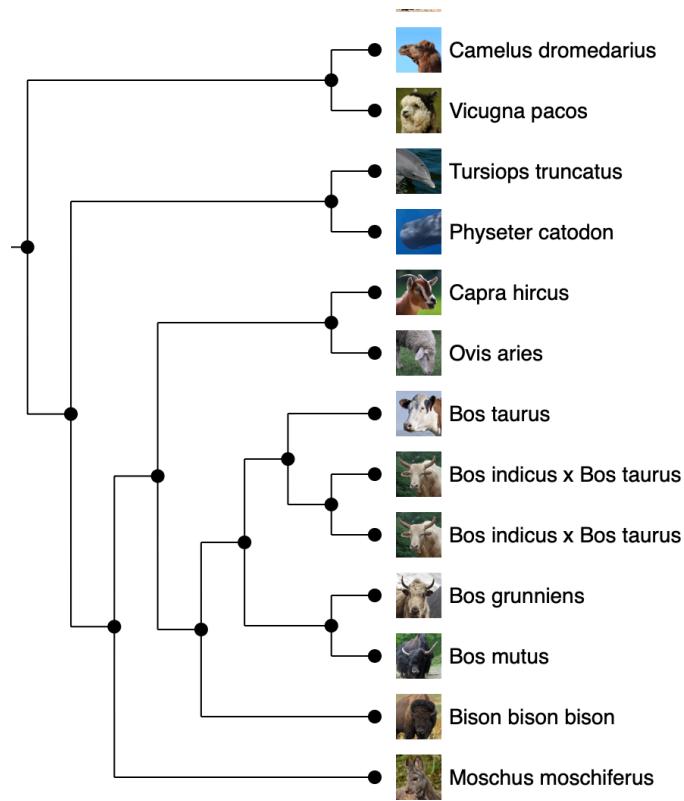
## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

# Sources of Genes Annotations



<https://www.gencodegenes.org>



<https://www.ensembl.org>



**March 6, 2020**  
**RefSeq Release 99 is available for FTP**

This release includes:

Proteins: 167,278,920  
Transcripts: 29,869,155  
Organisms: 99,842  
Available at: <ftp://ftp.ncbi.nlm.nih.gov/refseq/release/>  
Documentation: [Release Notes](#)

<https://www.ncbi.nlm.nih.gov/refseq/>

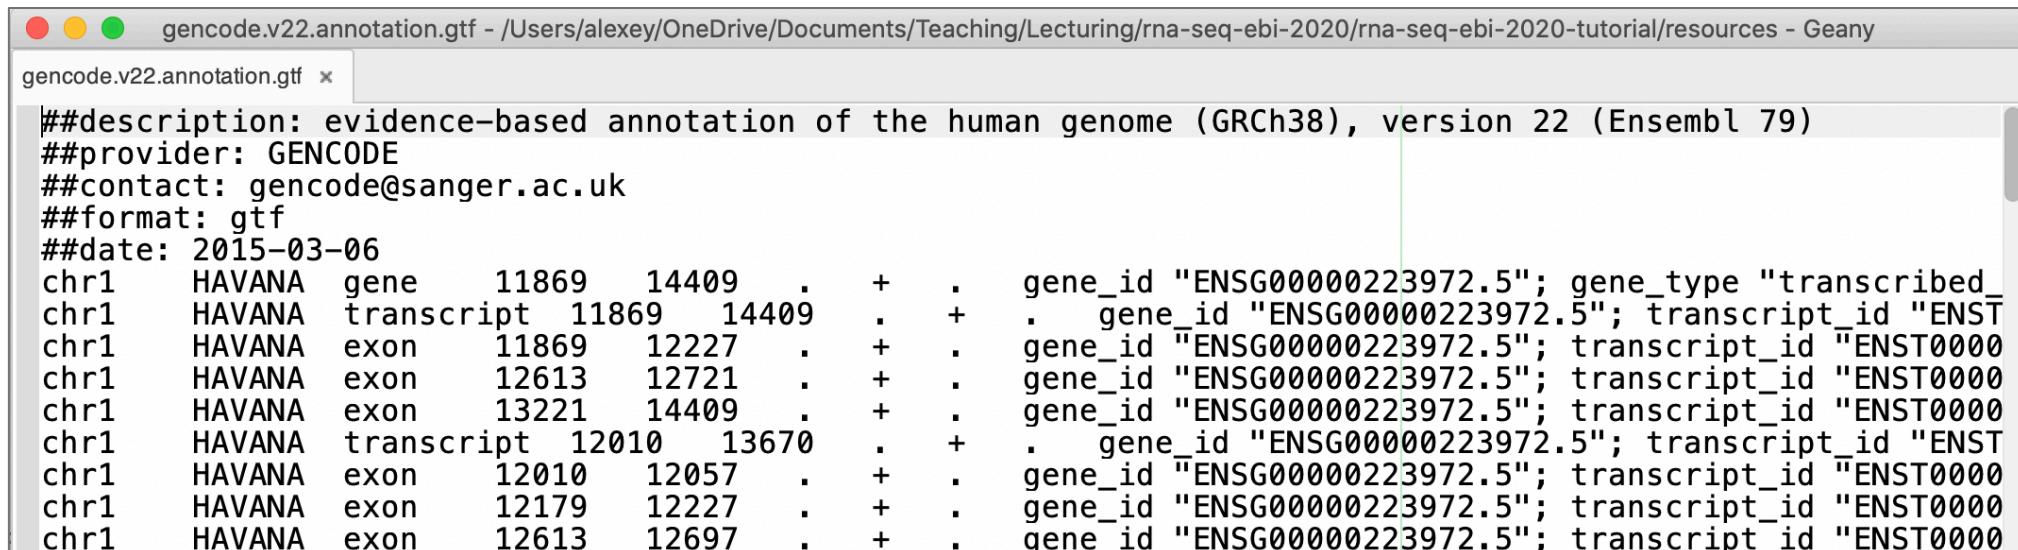
# GFF / GTF: file format for genes annotations

<https://www.ensembl.org/info/website/upload/gff.html>

<http://genome.ucsc.edu/FAQ/FAQformat.html#format3>

## GFF format

GFF (General Feature Format) lines are based on the Sanger [GFF2 specification](#). GFF lines have nine required fields that *must* be tab-separated. If the fields are separated by spaces instead of tabs, the track will not display correctly. For more information on GFF format, refer to Sanger's [GFF page](#).



```
##description: evidence-based annotation of the human genome (GRCh38), version 22 (Ensembl 79)
##provider: GENCODE
##contact: gencode@sanger.ac.uk
##format: gtf
##date: 2015-03-06
chr1 HAVANA gene 11869 14409 . + . gene_id "ENSG00000223972.5"; gene_type "transcribed_
chr1 HAVANA transcript 11869 14409 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST
chr1 HAVANA exon 11869 12227 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST0000
chr1 HAVANA exon 12613 12721 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST0000
chr1 HAVANA exon 13221 14409 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST0000
chr1 HAVANA transcript 12010 13670 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST
chr1 HAVANA exon 12010 12057 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST0000
chr1 HAVANA exon 12179 12227 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST0000
chr1 HAVANA exon 12613 12697 . + . gene_id "ENSG00000223972.5"; transcript_id "ENST0000
```

GTF files can be conveniently read into R data-frame using `readGFF()` function from `rtracklayer` package  
(see example in the practical session)

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

# Design formulas in R models

The “**design formulas**” are text strings, used to describe the required analysis for **edgeR** and **DESeq2**.

For simplicity, we will illustrate some “**design formulas**” using **lm()** function as an example.

Exactly the same “**formulas**” could be used for **edgeR** and **DESeq2**.

A simple design: Is the mean tumour size different between two groups ?

```
# Toy data
size <- c(1.35, 1.42, 1.75, 1.14, 2.67, 2.54, 3.98, 2.27)
group <- c("A","A","A","A","B","B","B","B")

# Combine to a data.frame
experiment.df <- data.frame(size, group)
str(experiment.df)

'data.frame': 8 obs. of 2 variables:
$ size : num 1.35 1.42 1.75 1.14 2.67 2.54 3.98 2.27
$ group: Factor w/ 2 levels "A","B": 1 1 1 1 2 2 2 2
```

	size	group
1	1.35	A
2	1.42	A
3	1.75	A
4	1.14	A
5	2.67	B
6	2.54	B
7	3.98	B
8	2.27	B

Note that the “group” is a Factor with base level “A”.  
This level will be used for the intercept (also called “reference” level).

# Is a mean tumour size different between two groups ?

Fit a linear model using "formula" to specify the model's "design"

**Intercept** =  
the mean size of tumours  
in **Group A**

Design formula  
**size ~ group**

The **Intercept** = size of tumours  
In **Group A** is significantly  
different from 0

```
> fit <- lm(size ~ group, data=experiment.df)
> coefficients(summary(fit))
   Estimate Std. Error t value Pr(>|t|) 
(Intercept)  1.415    0.2837986  4.98593 0.002486972
groupB       1.450    0.4013519  3.61279 0.011193446
```

The difference between  
the mean size of tumours  
in **Group B** and the **intercept**

The difference between  
**Group B** and the **intercept**  
is significantly different from 0

```

> fit <- lm(size ~ group, data=experiment.df)
> coefficients(summary(fit))
            Estimate Std. Error t value    Pr(>|t|)    
(Intercept)  1.415     0.2837986  4.98593 0.002486972
groupB       1.450     0.4013519  3.61279 0.011193446

```

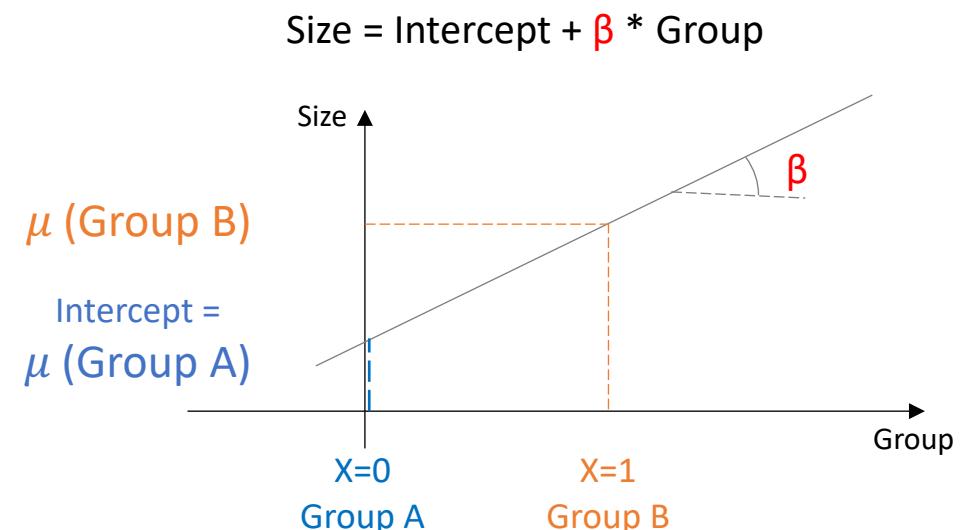
## Behind the scene

- 1) Coding the groups as 0 and 1 and (making “design matrix”)
- 2) Fitting an Intercept (default: the group with the base level in the factor)
- 3) Calculating the slope ( $\beta$ )

```

> model.matrix(~ group, data=experiment.df)
(Intercept) groupB
1           1     0
2           1     0
3           1     0
4           1     0
5           1     1
6           1     1
7           1     1
8           1     1

```



# Setting Intercept in R linear models

Default:

implicit modeling with intercept

“ $Y \sim Group$ ” in fact means “ $Y \sim Intercept + Group$ ”

```
> fit <- lm(size ~ group, data=experiment.df)
> coefficients(summary(fit))
            Estimate Std. Error t value    Pr(>|t|)
(Intercept)  1.415   0.2837986 4.98593 0.002486972
groupB       1.450   0.4013519 3.61279 0.011193446
```

Overriding the default:

explicit modeling without intercept

$Y \sim 0 + Group$

```
> fit <- lm(size ~ 0 + group, data=experiment.df)
> coefficients(summary(fit))
            Estimate Std. Error t value    Pr(>|t|)
groupA     1.415   0.2837986 4.98593 2.486972e-03
groupB     2.865   0.2837986 10.09519 5.487252e-05
```

```
> model.matrix(~ group, data=experiment.df)
(Intercept) groupB
1           1     0
2           1     0
3           1     0
4           1     0
5           1     1
6           1     1
7           1     1
8           1     1
```

```
> model.matrix(~ 0 + group, data=experiment.df)
groupA groupB
1       1     0
2       1     0
3       1     0
4       1     0
5       0     1
6       0     1
7       0     1
8       0     1
```

# Setting Intercept in R linear models

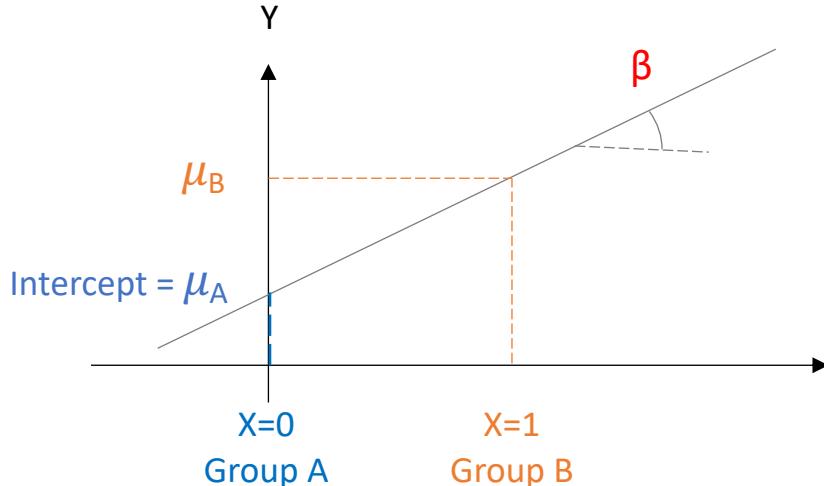
Understand what you are asking for ...

Don't override the defaults without asking a statistician ....

**Default:**

implicit modeling with intercept

"*Y ~ Group*" in fact means "*Y ~ Intercept + Group*"



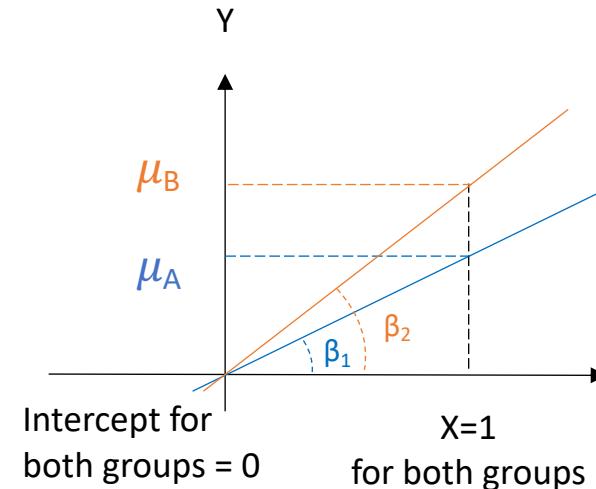
Significance of  $\beta \neq 0$

Look at the difference of mean values  
between Group A and Group B

**Overriding the default:**

explicit modeling without intercept

$Y \sim 0 + Group$



Significance of  $\beta_1 \neq 0$  or  $\beta_2 \neq 0$

Look at means in both groups:  
whether each is different from zero

Hmm...  
This is interesting ...  
But this is **NOT** what we  
have been looking for!

## Design with a confounder (batch effect)

What if we measured the tumour sizes over two days, and on the second day we had a different calibration, so all the measurements on the second day somehow went higher?

```
# Toy data
size <- c(1.42, 1.35, 1.95, 1.87, 2.67, 2.54, 3.98, 3.27)
group <- c("A", "A", "B", "B", "A", "A", "B", "B")
day <- c("Mon", "Mon", "Mon", "Mon", "Tue", "Tue", "Tue", "Tue")

# Combine to a data.frame
experiment.df <- data.frame(size, group, day)
str(experiment.df)

'data.frame': 8 obs. of 3 variables:
 $ size : num 1.42 1.35 1.95 1.87 2.67 2.54 3.98 3.27
 $ group: Factor w/ 2 levels "A","B": 1 1 2 2 1 1 2 2
 $ day   : Factor w/ 2 levels "Mon","Tue": 1 1 1 1 2 2 2 2
```

	size	group	day
1	1.42	A	Mon
2	1.35	A	Mon
3	1.95	B	Mon
4	1.87	B	Mon
5	2.67	A	Tue
6	2.54	A	Tue
7	3.98	B	Tue
8	3.27	B	Tue

We can see clearly the differences between groups within each day ...  
However, how to combine the data over both days (batches) ?

# Including batch correction in the design formula

## Formula without batch correction

size ~ group

No significant association between Group and Size

```
> fit <- lm(size ~ group, data=experiment.df)
> coefficients(summary(fit))
   Estimate Std. Error t value Pr(>|t|)    
(Intercept) 1.9950    0.4423317 4.51019 0.00405967
groupB      0.7725    0.6255514 1.23491 0.26302689
```

```
> model.matrix(~ group, data=experiment.df)
```

	(Intercept)	groupB
1	1	0
2	1	0
3	1	1
4	1	1
5	1	0
6	1	0
7	1	1
8	1	1

## Formula with batch correction

size ~ day + group

Detected an association between Group and Size

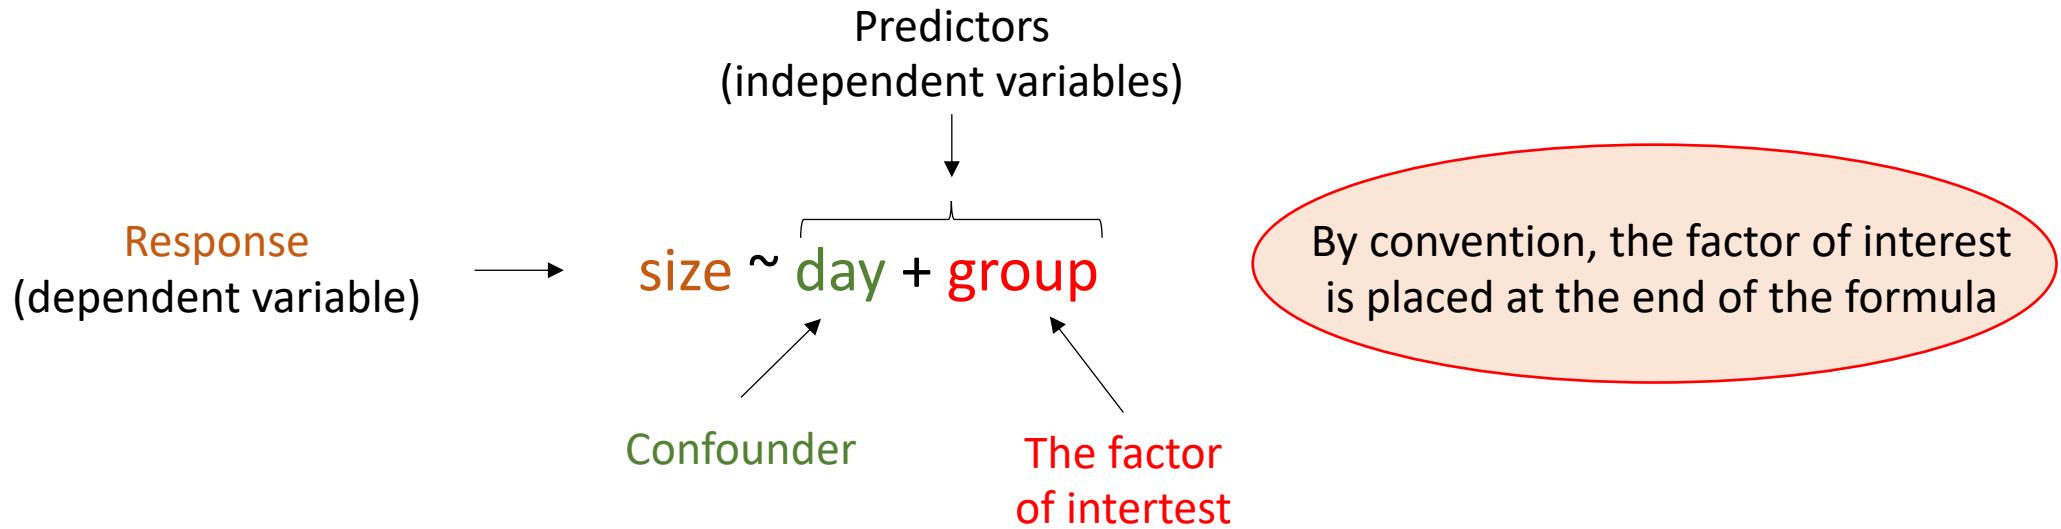
```
> fit <- lm(size ~ day + group, data=experiment.df)
> coefficients(summary(fit))
   Estimate Std. Error t value Pr(>|t|)    
(Intercept) 1.26125   0.1707328 7.387274 0.0007148004
dayTue      1.46750   0.1971453 7.443750 0.0006899725
groupB      0.77250   0.1971453 3.918431 0.0111996428
```

```
> model.matrix(~ day + group, data=experiment.df)
```

	(Intercept)	dayTue	groupB
1	1	0	0
2	1	0	0
3	1	0	1
4	1	0	1
5	1	1	0
6	1	1	0
7	1	1	1
8	1	1	1

# Including covariates / confounders in the design formula

## The previous example in more detail



## Other examples

Is a treatment significantly associated with change in blood pressure, controlling for age and ethnicity?

**blood\_pressure\_change ~ age + ethnicity + treatment**

Is a gene differentially expressed in different types of tumour, controlling for batch and sex?

**tumour\_type ~ batch + sex + gene**

# Design with an interaction

May one factor change the response to another ?

Consider an imaginary experiment:

We start treatment of breast tumours in mice when they reached a certain size (e.g. 1cm).

Then we record the size after one week of treatment, and information about the type of tumour (ER-pos or ER-neg) and about the type of treatment (Tamoxifen or Placebo).

```
# Make a toy dataset
size <- c(0.35, 0.42, 1.3, 1.4, 1.3, 1.2, 1.1, 1.4)
er <- c("Pos", "Pos", "Pos", "Pos", "Neg", "Neg", "Neg", "Neg")
tamoxifen <- c("Yes", "Yes", "No", "No", "Yes", "Yes", "No", "No")

# Combine to a data.frame
experiment.df <- data.frame(size, er, tamoxifen)
str(experiment.df)

'data.frame': 8 obs. of 3 variables:
 $ size     : num  0.35 0.42 1.3 1.4 1.3 1.2 1.1 1.4
 $ er       : Factor w/ 2 levels "Neg","Pos": 2 2 2 2 1 1 1 1
 $ tamoxifen: Factor w/ 2 levels "No","Yes": 2 2 1 1 2 2 1 1
```

	size	er	tamoxifen
1	0.35	Pos	Yes
2	0.42	Pos	Yes
3	1.30	Pos	No
4	1.40	Pos	No
5	1.30	Neg	Yes
6	1.20	Neg	Yes
7	1.10	Neg	No
8	1.40	Neg	No

We expect that ER status modifies the response to tamoxifen:

- ER-positive tumours should shrink upon Tamoxifen treatment while
  - ER-negative tumours should keep growing

# Including interaction in the design formula

## Formula without interaction

size ~ er + tamoxifen

No significant  
**Size** change associated  
with **Tamoxifen** or **ER**  
alone

```
> fit <- lm(size ~ er + tamoxifen, data=experiment.df)
> coefficients(summary(fit))
   Estimate Std. Error t value Pr(>|t|)
(Intercept) 1.49125  0.1980649  7.529099 0.0006543625
erPos       -0.38250  0.2287056 -1.672456 0.1552945113
tamoxifenYes -0.48250  0.2287056 -2.109699 0.0886557090
```

```
> model.matrix(~ er + tamoxifen, data=experiment.df)
(Intercept) erPos tamoxifenYes
1           1     1          1
2           1     1          1
3           1     1          0
4           1     1          0
5           1     0          1
6           1     0          1
7           1     0          0
8           1     0          0
```

## Formula with interaction

size ~ er \* tamoxifen

Significant  
**Size** change because  
of interaction between  
**Tamoxifen** and **ER**

```
> fit <- lm(size ~ er * tamoxifen, data=experiment.df)
> coefficients(summary(fit))
   Estimate Std. Error t value Pr(>|t|)
(Intercept) 1.250  0.08474226 14.7506103 0.0001229477
erPos       0.100  0.11984365  0.8344205 0.4509786113
tamoxifenYes 0.000  0.11984365  0.0000000 1.0000000000
erPos:tamoxifenYes -0.965 0.16948451 -5.6937356 0.0047004777
```

```
> model.matrix(~ er * tamoxifen, data=experiment.df)
(Intercept) erPos tamoxifenYes erPos:tamoxifenYes
1           1     1          1          1
2           1     1          1          1
3           1     1          0          0
4           1     1          0          0
5           1     0          1          1
6           1     0          0          1
7           1     0          0          0
8           1     0          0          0
```

## Another example of interpreting design formulas

RNA-seq was used to measure genes expression in three squamous cell carcinoma patients. Tumour and paired Normal tissue was available for each patient. Three types of model design was considered in the study:

Model	Interpretation	Genes detected
Patient	Baseline patient differences	
Patient + tissue	Consistent tumour differences	1276
Patient × tissue	Patient-specific tumour differences	202

McCarthy et al, 2012: Differential expression analysis of multifactor RNA-Seq experiments ...

<https://academic.oup.com/nar/article/40/10/4288/2411520>

# The design formula summary

- At the left side of the formula: The “Response” variable = “Dependent” variable (e.g. **Size** in the above examples)
- At the right side of the formula: “Predictors” = “Independent” variables (e.g. **Day**, **Group** etc in some above examples)
- The “predictors” include:
  - The variable of interest (e.g. **Group** in some above examples)
  - Confounders / covariates, which effect should be “controlled for” (e.g. **Day** in the batch effect example)
- By convention, most of the packages expect the variable of interest at the end of the formula
- Be aware about the reference level within the variable of interest (by default, the “base” level in the factor)
- Be aware about complex designs, such as designs with interactions, designs without intercepts. Other complex analyses may use user-defined contrasts etc. Ask a statistician if you consider using complex designs.

---

For simplicity, we illustrated design formulas using ***lm()*** function. However, many other R functions and packages, including ***glm()***, ***edgeR*** and ***DESeq2*** use the same way of specifying the design of analysis.

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

# Internal representation of data in DESeq2 and edgeR

## DESeq2 Data-Set

**Design slot**  
 $y \sim a + b$

**rowData slot**  
for genes data  
(e.g. ID, genomic coordinates etc)

**colData slot**  
for samples data  
(e.g. ID, diagnosis etc)

samples

genes

**Assay slot(s)**  
for Counts  
and other data

A modified ***Summarized Experiment***

## edgeR DGEList

A list containing 3 synchronized objects:

- **Samples** data frame
- **Genes** data frame
- **Counts** matrix

Although is not implemented  
as ***Summarized Experiment***  
provides similar functionality

# Import functions and packages

## *readDGE function in edgeR*

Makes a single matrix from multiple **text files of arbitrary format**, as long as the file contains columns with genes names and counts. Allows to add samples information etc. See example in the practical session.

<https://master.bioconductor.org/packages/release/workflows/vignettes/RNAseq123/inst/doc/limmaWorkflow.html>

## *DESeqDataSetFrom -Matrix, -HTSeqCount and -Tximport functions in DESeq2*

Allow convenient import from matrix (**Rsubread**), **HTSeq** counts and from **Tximport** package respectively.

<https://www.bioconductor.org/packages-devel/workflows/vignettes/rnaseqGene/inst/doc/rnaseqGene.html>

## **Tximport package**

Facilitates import from **Salmon**, **Kallisto** and some other tools to **DESeq2** or **edgeR**. May summarize transcripts to genes.

<https://bioconductor.org/packages-devel/bioc/vignettes/tximport/inst/doc/tximport.html>

## **Tximeta package**

Currently facilitates import from **Salmon** to **DESeq2**, May summarize transcripts to genes.

Adds meta-data, including genes annotation in Genomic Ranges format.

<https://bioconductor.org/packages/release/bioc/vignettes/tximeta/inst/doc/tximeta.html>

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

**Normalizing by library size**  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering

Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

## Visualizing results

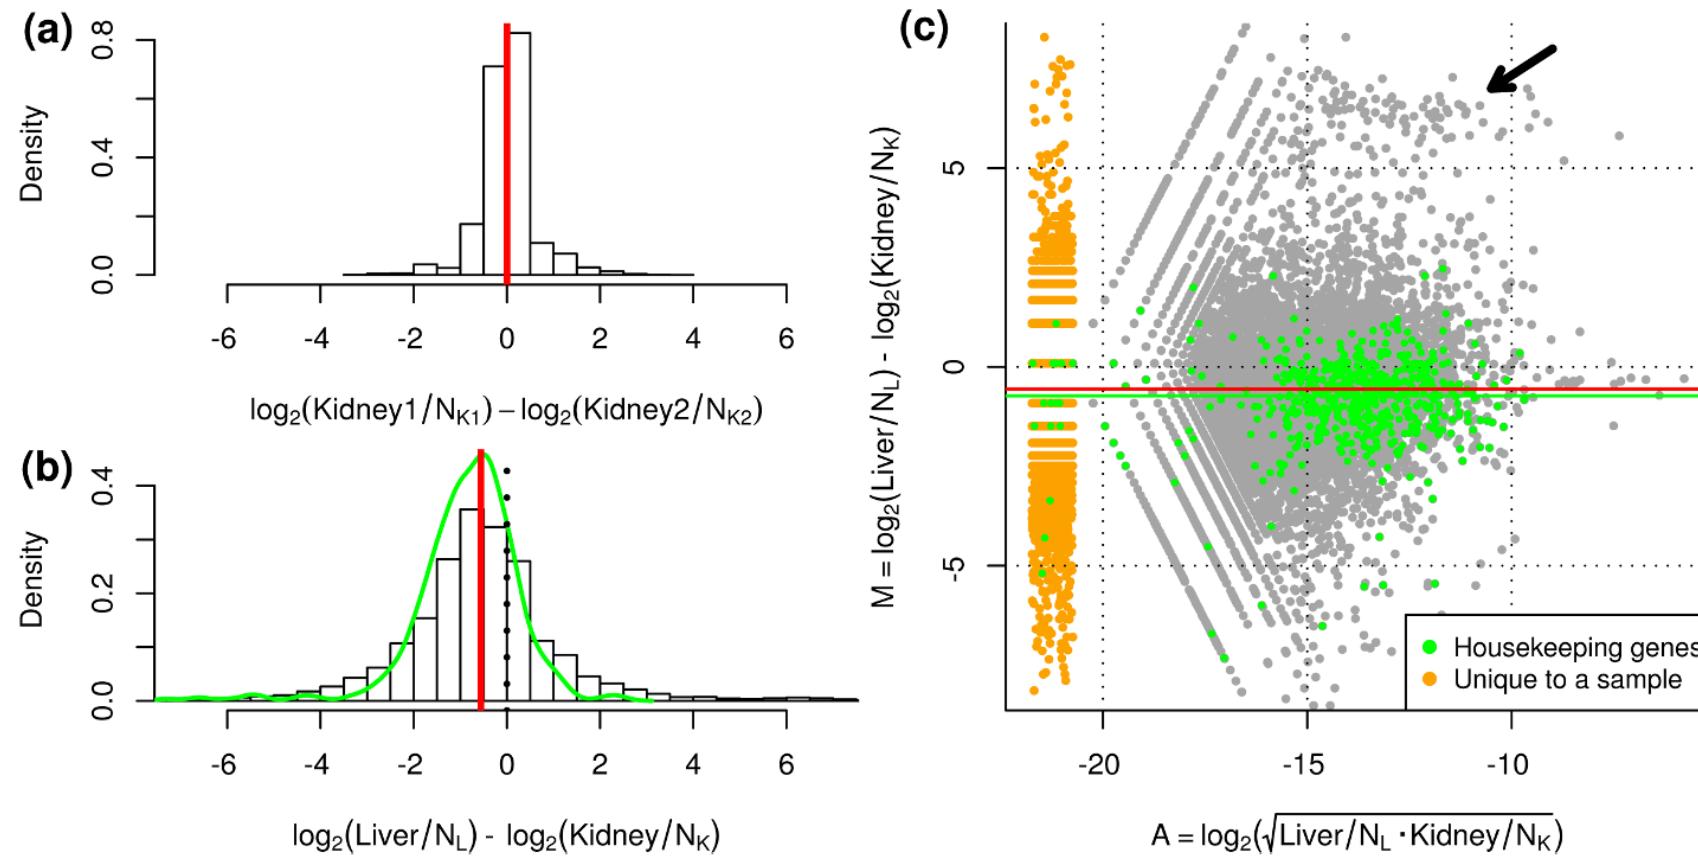
MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

# Normalizing by library size

Initially, a naïve normalization was suggested by the library size itself. However, later it was observed that, with such naïve normalization, a strong over-expression of some genes might be mis-interpreted as down-regulation of all the other genes:



# TMM and Median-of-Ratios

To avoid such undue influence, **edgeR** excludes extremely expressed and extremely changed genes when calculating the normalization factors. The method is called “trimmed mean of M values” (**TMM**)

Robinson and Oshlack 2010: A scaling normalization method for differential expression analysis of RNA-seq data  
<http://genomebiology.com/2010/11/3/R25>

For the same purpose, to avoid the undue influence of extremely changed genes, **DESeq2** takes **Median of the Genes Ratios** to estimate the size factors. The statistical properties of the **median** negate the effect of extremely changed genes.

Anders and Huber 2010: Differential expression analysis for sequence count data  
<http://genomebiology.com/2010/11/10/R106>

## TMM and Median of the Gene Ratios do not normalize by gene length

In contrast to some “normalized” units used for RNA-seq counts representation (FPKM or TPM), **TMM** or **Median of the Gene Ratios** do not account for the gene/transcript length.

This implicitly assumes that the gene/transcript length does not change between the studied conditions. Although such assumption looks reasonable for Differential **Genes** Expression analysis on short-read data, changes in size of used transcripts may be incorporated in the modelling later.

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
**Exploring the source data**

- Variance-stabilizing transformations
- PCA and Hierarchical clustering

Dispersion estimates

- Need for “borrowing” data
- Empirical adaptive estimates

Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

# Variance-stabilizing transformations to explore source data

It is a common practice to perform unsupervised clustering of samples using all (or most variable) genes before the Differential Gene Expression analysis. Such exploration may suggest gross outliers, to exclude from analysis.

Also, it shows whether the studied groups are well separated in the gene expression space.

The most common methods for samples clustering include **Principal Component Analysis (PCA)** and **Hierarchical Clustering**.

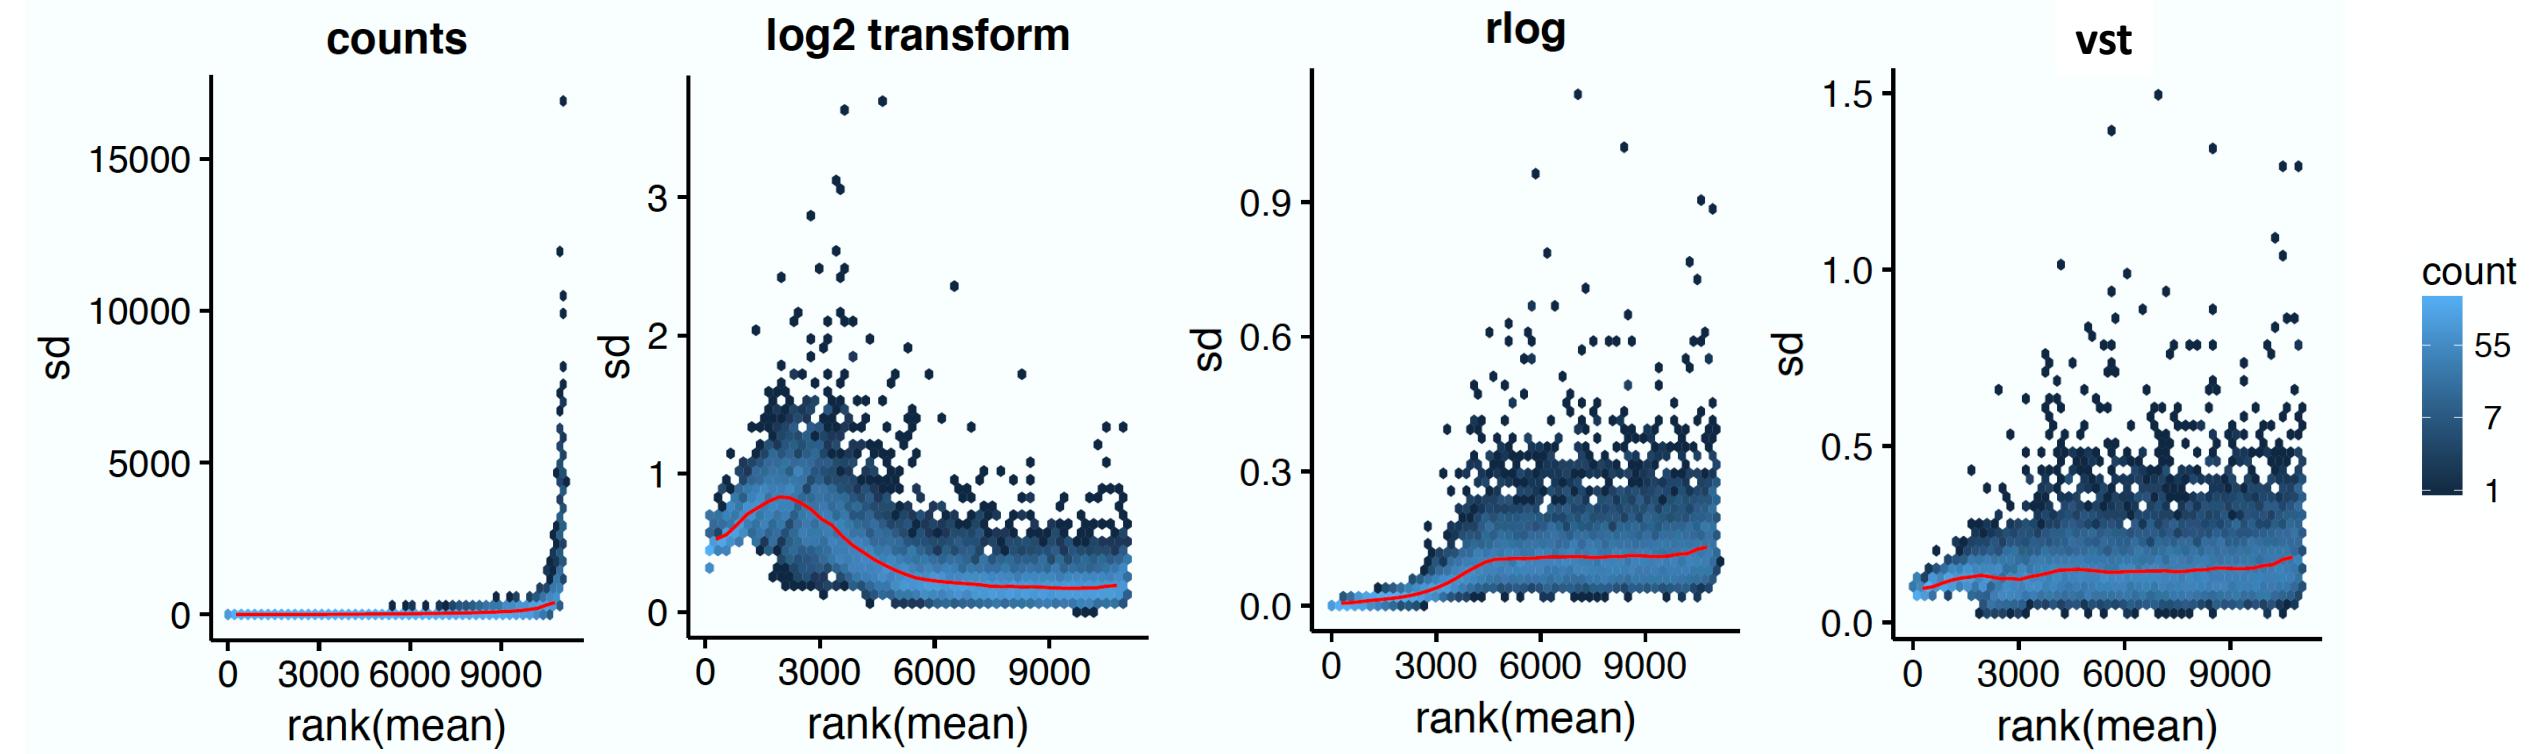
However, it has been empirically observed that variance in RNA-seq counts is higher in the highly-expressed genes (**heteroskedasticity** of RNA-seq counts). Thus, **PCA** and **Hierarchical Clustering** might be dominated by the most expressed genes if the data are not transformed to make variance similar between the genes (make transformed data **homoscedastic**).

Both **edgeR** and **DESeq2** provide transformations and plotting functions to facilitate the exploratory analysis:

	<b>edgeR</b>	<b>DESeq2</b>
Transformation(s)	Log(counts per million)	VST (Variance-Stabilizing transformation) rlog (regularized log-transformation)
Plotting function	plotMDS*()	plotPCA()

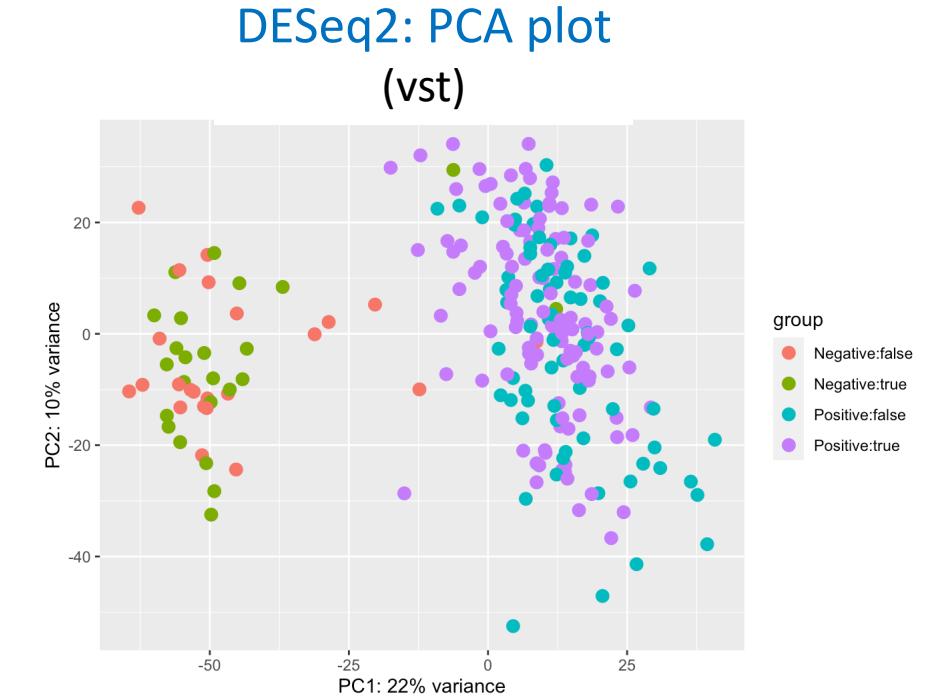
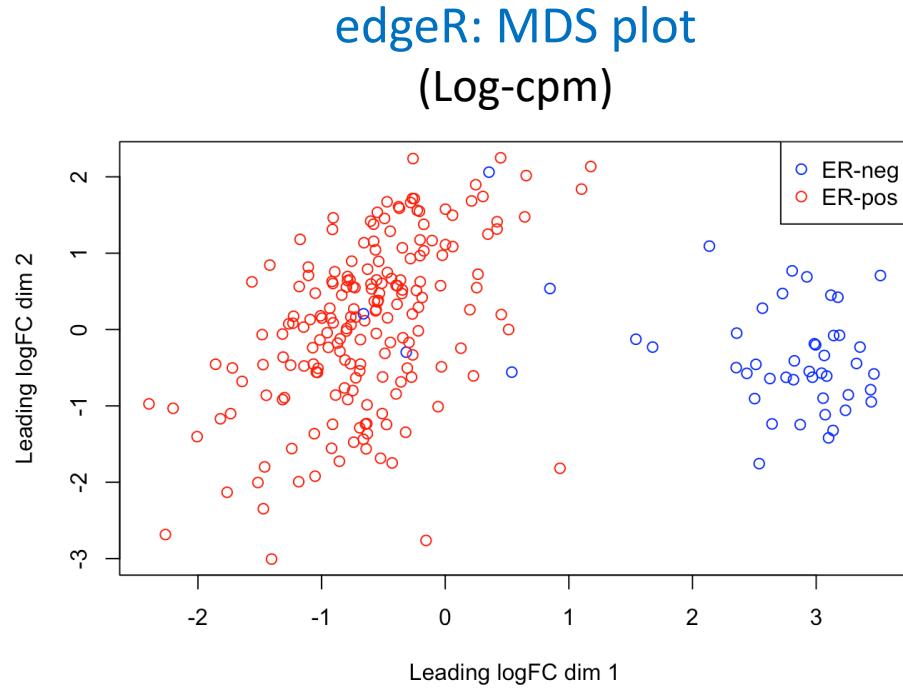
\* **MDS** stands for **Multi-Dimensional Scaling**: this is a procedure very similar to PCA

## Variance-stabilizing transformations to explore source data



These transformations are used only to explore source data  
They are NOT used during the Differential Expression Analysis

# Examples of PCA and MDS plots



These plots will be generated during the practical session (along with Hierarchical Clustering and Heatmap lots).  
The plots don't suggest any gross outliers to exclude. Also they suggest that some of the studied groups are clearly separated in the gene expression space.

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering

### Dispersion estimates

- Need for “borrowing” data
- Empirical adaptive estimates

Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

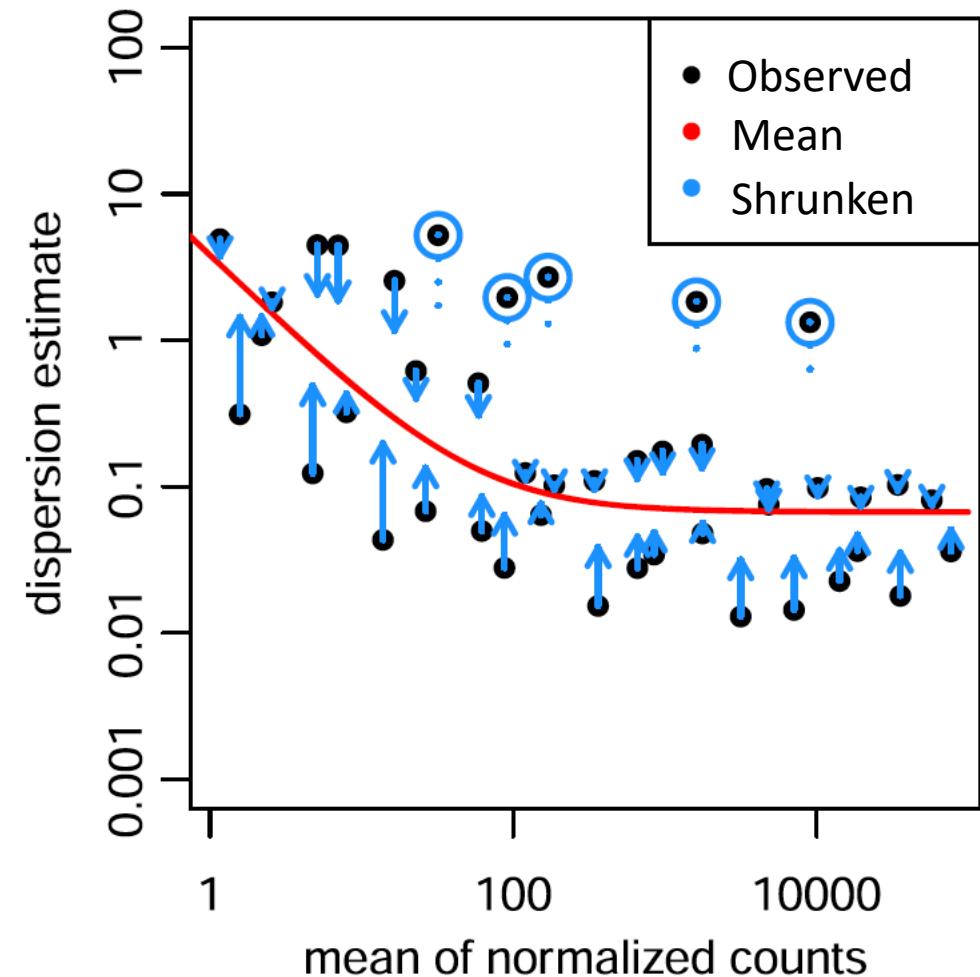
Comparison of methods  
Comparison of results

# Dispersion estimation and adjustment

If dispersion for each single gene can not be accurately estimated because of a small number of samples (e.g. less than 10 replicates) then the data from other genes will be “borrowed”.

Simplified description of the procedure applied by DESeq2 :

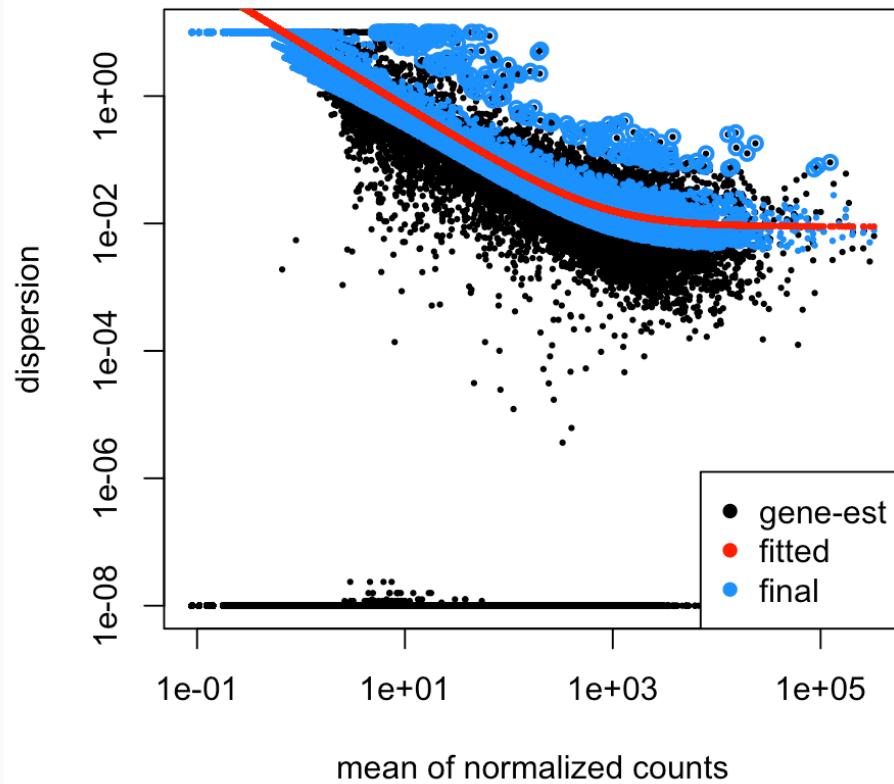
1. **Observed** dispersions (●) are used to estimate **Mean** dispersions (●) for each level of expression.
2. Depending on the accuracy of the **Observed** dispersions they may be “**Shrunken**” (●) toward the **Mean** estimates. The more accurate is the observed dispersion, the less “shrinkage” will be applied.
3. If the **Observed** dispersion extremely deviates from **Mean** (outliers encircled in blue) it does not shrink.



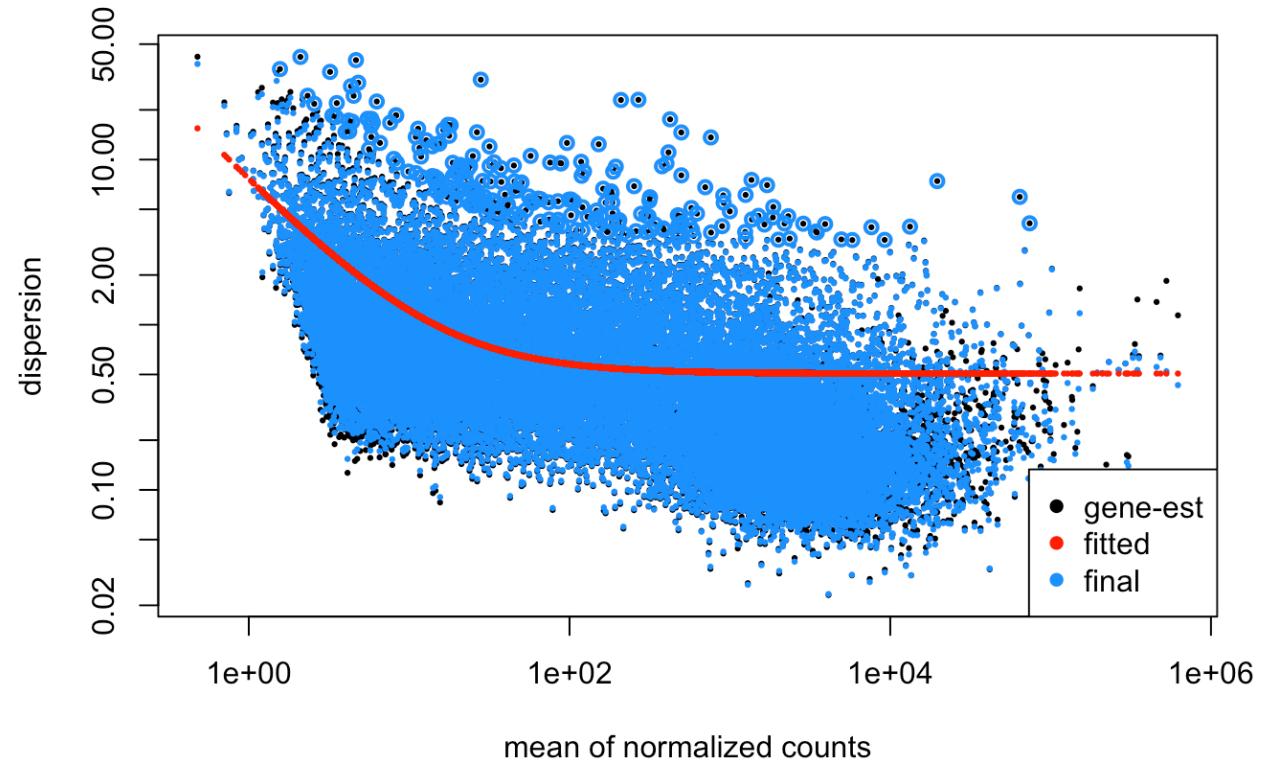
Love et al 2014: Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2

<https://genomebiology.biomedcentral.com/articles/10.1186/s13059-014-0550-8>

# Dispersion estimation and adjustment



Dataset with small number of samples.  
Many dispersions are shrunken toward the mean.



Dataset with large number of samples.  
Most dispersions are not shrunken toward the mean.

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates

## Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

## Default thresholds in DESeq2

It is a commonly accepted assumption that in most of experiments only a minority of genes change their expression. Thus, usually, the proportion of suggested Differentially Expressed Genes should not exceed 20%. However, in our practical session the default **DESeq2** settings would suggest that about 60% of all genes are differentially expressed ...

This is obviously an absurd result: the problem is in the default thresholds applied by DESeq2.

By default **DESeq2** performs testing for **any fold change at FDR < 0.1**.

Such settings reflect the time, when RNA-seq was prohibitively expensive and experiments often included less than 10 samples.

Our practical session analyses more than 200 samples. For such datasets, **DESeq2** allows to change the analysis thresholds. Thus, when we consider only genes with **at least 2-fold change at FDR < 0.01**, the proportion of suggested DEGs got below 10%.

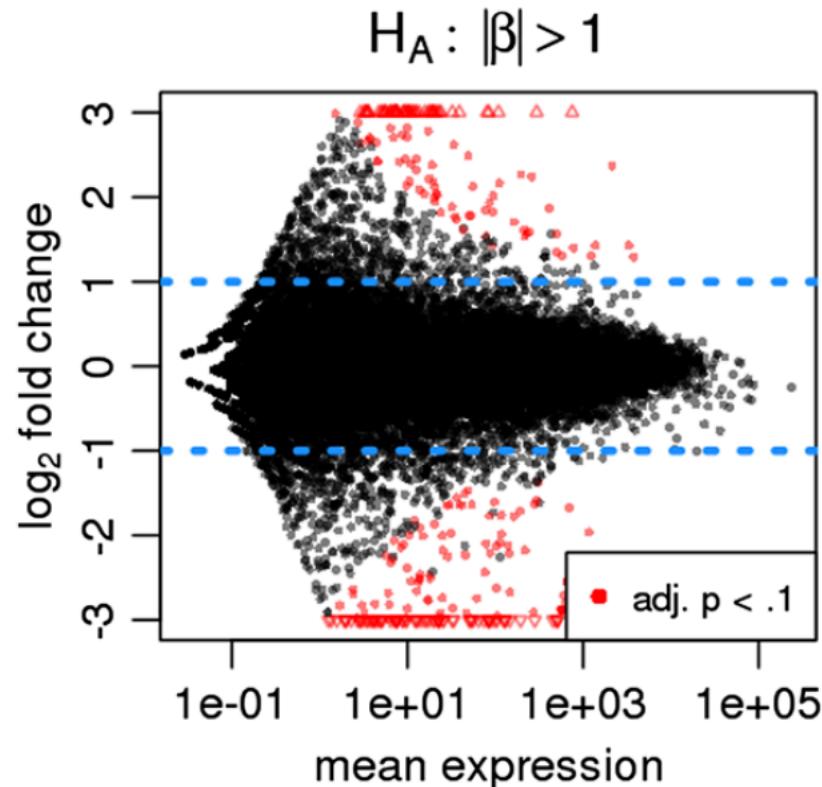
**edgeR** also allows testing for non-zero Fold-Change (see practical session for examples).

Importantly, testing against non-zero FC threshold is not the same as filtering by FC of the results obtained with default settings.

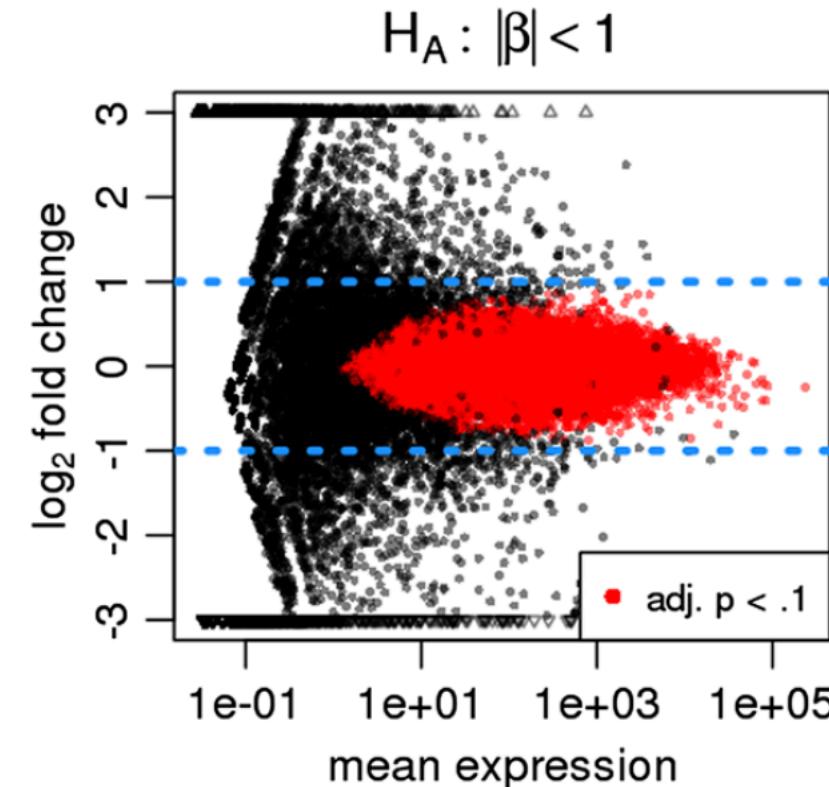
# Testing against non-zero FC thresholds

MA-plots show Fold-Change (Y-axis) against Mean expression (X-axis) for individual genes; red shows significance

Genes with at least 2-fold change



Genes with less than 2-fold change



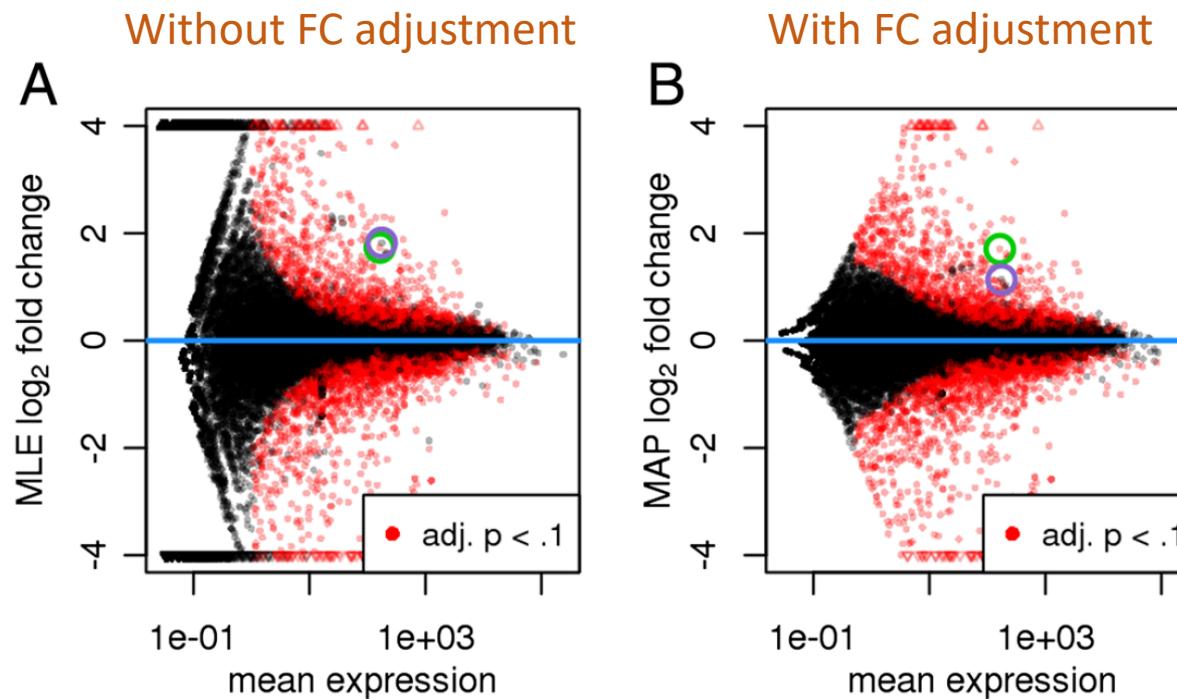
Love et al 2014: Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2

<https://genomebiology.biomedcentral.com/articles/10.1186/s13059-014-0550-8>

# Additional features in DESeq2

## Fold-Change (FC) adjustment in low-expressed genes

Noise may simulate high Fold-Changes in low-expressed genes. To avoid this artificially inflated “Changes” **DESeq2** uses an empirical algorithm that “shrinks” fold change in the genes with low expression.

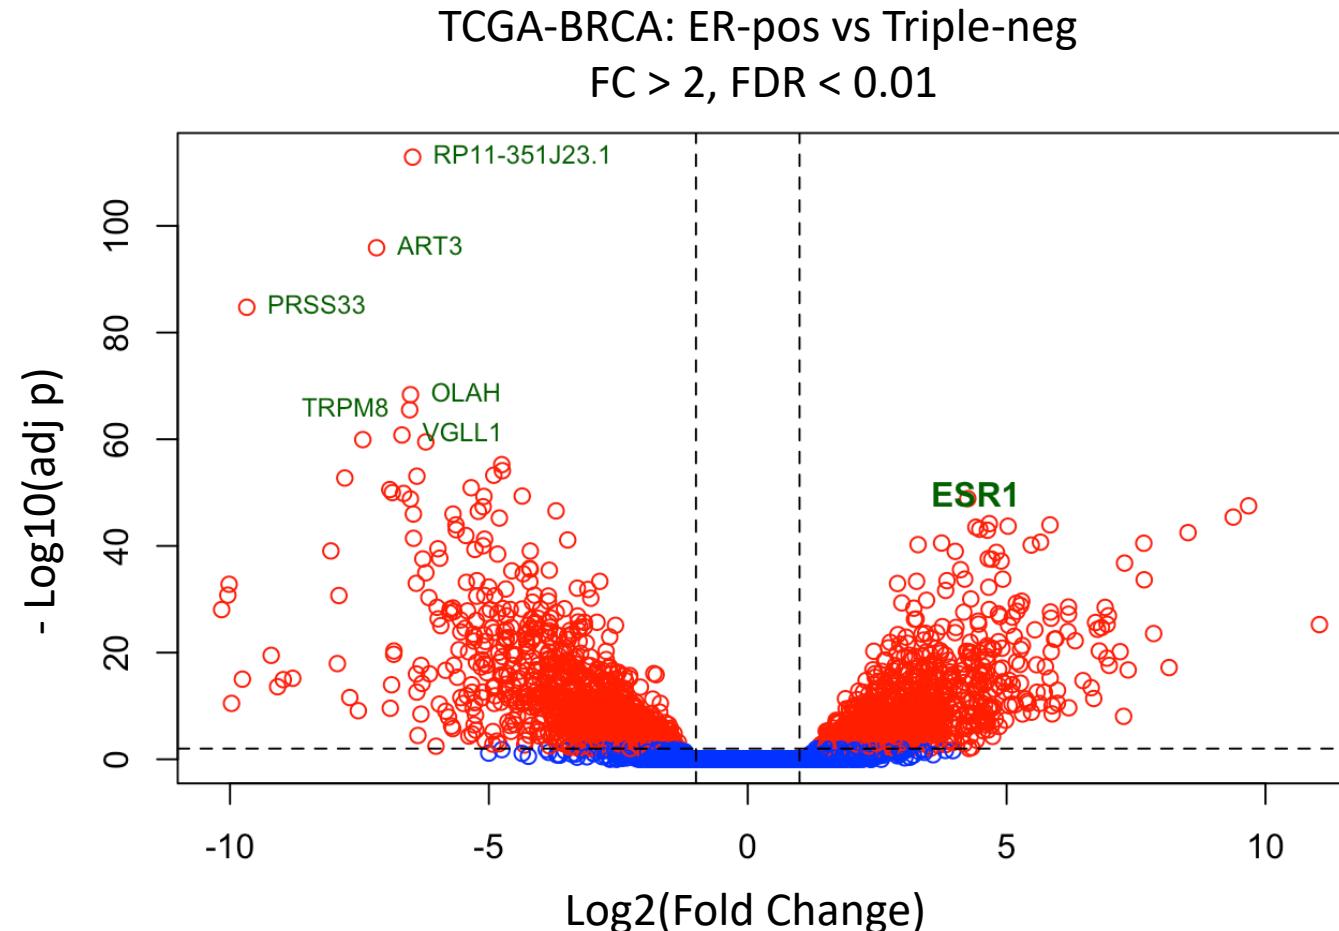


## Removal of low-expressed genes before calculating FDR

DESeq2 automatically removes genes with low expression, applying adaptive threshold that maximises number of genes passing FDR. For the genes filtered out at this stage, NA is placed in the **p-adjusted** column.

# Example of Volcano plot

**Volcano-plot** shows Significance (Y-axis) against Fold-Change (X-axis) for each gene.  
Genes could be coloured according to FC and Significance thresholds



Result from the practical session

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial Distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results

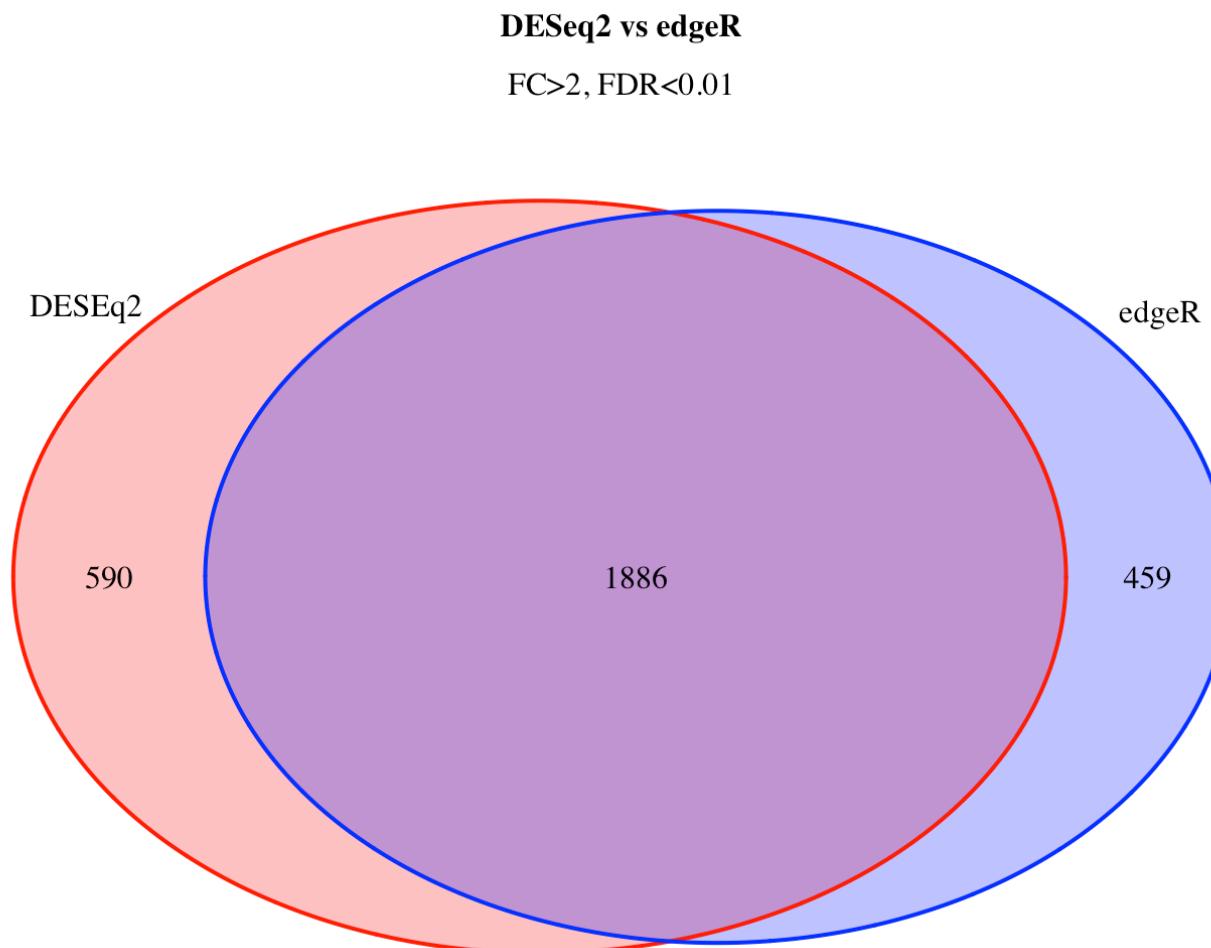
# Statistical features of edgeR and DESeq2

Step	edgeR	DESeq2
Normalizing by library size	Trimmed Mean of M-values	Median of the Genes' Ratios
Distribution	Negative Binomial	Negative Binomial
Dispersion for a gene	An empirical custom procedure accounting for (i) dispersion over all genes, (ii) dispersion in the genes with similar expression and (iii) dispersion observed in the gene	An empirical custom procedure accounting for (i) dispersion in the genes with similar expression and (iii) dispersion observed in the gene
GLM	Log-link optimized algorithm for convergence	Log-link
Significance test	Log-likelihood ratio	Wald test
Multiple testing correction	FDR	FDR

# Accessory features of edgeR and DESeq2

Step	edgeR	DESeq2
Internal data format	Customized list: DGEList	Modified Summarized Experiment
Data import	Data import function from a set of text files containing columns for <b>gene-ID</b> and <b>gene-Counts</b>	<ul style="list-style-type: none"> <li>- Advanced import options are provided by <b>tximport</b> and <b>tximeta</b> packages</li> <li>- DESeq2 includes data import functions from different upstream tools (tximport, HTSeq, matrix)</li> </ul>
Low-expressed genes	A function for filtering by low expression (applied before analysis)	<ul style="list-style-type: none"> <li>- Automatic exclusion of low-expressed genes from multiple testing</li> <li>- Adjusting Fold Change for low-expressed genes</li> </ul>
Testing against non-zero fold change (FC)	Yes	Yes, with an opportunity of testing for FC below or above the Threshold
Detecting counts outliers	No	Count outliers are detected and taken into account when calculating FC, p and when filtering genes
Data exploration functions	Log(counts per million) MDS plot	Variance-Stabilising transformation (VST) Regularized log-transformation (rlog) PCA plot

## Example of DEGs detected by DESeq2 and edgeR



Result from the practical session

# Differential Gene Expression

## Introduction

Short and Long reads: Genes and Transcripts  
Upstream and down-stream applications

## Statistical summary

Recap of “standard” approaches  
Problems with RNA-seq counts  
Overdispersion: Negative Binomial distribution

## Counts

Overview and software  
Units: raw counts, (R)FPKM, TPM

## Genes

GENCODE, Ensembl, Refseq  
GTF file format

## Design formula

A simple design  
Accounting for covariates e.g. batch effect  
Advanced designs: interactions

## Data import

Summarized Experiment and DGEList  
Data import packages and functions

## Statistics in more details

Normalizing by library size  
Exploring the source data  
- Variance-stabilizing transformations  
- PCA and Hierarchical clustering  
Dispersion estimates  
- Need for “borrowing” data  
- Empirical adaptive estimates  
Total number of DEGs and thresholds selection

## Visualizing results

MA- and Volcano plots

## DESeq2 vs edgeR

Comparison of methods  
Comparison of results



## Selected references

### DESeq2

<https://www.bioconductor.org/packages/devel/workflows/vignettes/rnaseqGene/inst/doc/rnaseqGene.html>

Love *et al*, 2014: Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2

<https://genomebiology.biomedcentral.com/articles/10.1186/s13059-014-0550-8>

Anders and Huber, 2010: Differential expression analysis for sequence count data

<http://genomebiology.com/2010/11/10/R106>

### edgeR

<https://master.bioconductor.org/packages/release/workflows/vignettes/RNAseq123/inst/doc/limmaWorkflow.html>

Robinson and Smyth, 2007: Moderated statistical tests for assessing differences in tag abundance

<https://academic.oup.com/bioinformatics/article/23/21/2881/372869>

Robinson and Oshlack, 2010: A scaling normalization method for differential expression analysis of RNA-seq data

<http://genomebiology.com/2010/11/3/R25>

McCarthy *et al*, 2012: Differential expression analysis of multifactor RNA-Seq experiments with respect to biological variation

<https://academic.oup.com/nar/article/40/10/4288/2411520>

## Practical session

- Using open-access RNA-seq data for several hundred samples from **TCGA dataset** the tutorial will provide step-by step instructions on how to detect genes differentially expressed between Estrogen-Receptor-positive and Triple-negative breast cancers.
- During the practical session you will
  - Import **HTSeq counts** to **DESeq2** and **edgeR** data formats
  - Add information about the **samples** and **genes**
  - Remove consistently low-expressed **genes**
  - Perform **normalization** to account for library sizes
  - Explore source data using **PCA** and **MDS** plots, perform **Hierarchical Clustering** and make **Heatmap** plot to show clustering of samples and genes
  - Identify Differentially Expressed Genes with at least **2-fold difference at FDR < 0.01**
  - Explore plots of the **dispersion estimates and adjustments**
  - Explore **MA- and Volcano plots** for the Differentially Expressed Genes
- The equivalent analyses will be performed by **DESeq2** and **edgeR**, and the results will be compared between these packages